Evolution, Semantics, and Engineering of the F⋆ Verification System

May 19, 2019

Principal Investigator:
Cătălin Hriţcu ⟨catalin.hritcu@inria.fr⟩

Instittion: INRIA Paris

Introduction

F⋆ [2, 3, 17, 31] is a general-purpose functional programming language with effects aimed at program verification. It puts together the automation of an SMT-backed deductive verification tool with the expressive power of a proof assistant based on dependent types [21]. After verification, F* programs can be extracted to efficient OCaml, F#, C [28], WASM [27], or ASM code. The main ongoing use case of F* is building a verified, drop-in replacement for the whole HTTPS stack in Project Everest [9], including the underlying cryptographic primitives [35].

F⋆ is an open source project developed in the open and having an inclusive collaboration policy that has attracted contributors from Microsoft Research, Inria Paris, Nomadic Labs, University of Ljubljana, University of Edinburgh, Zen Protocol, Princeton University, MIT, etc. The code of F* is released under the permissive Apache License v2 and is developed on GitHub (at https://github.com/FStarLang/FStar).

While the F⋆ verification system shows great promise in practice, many challenging conceptual problems remain to be solved, many of which can directly inform the further evolution and design of the language. Moreover, many engineering challenges remain in order to build a truly usable verification system. This proposal promises to help address this by focusing on the following 5 main topics:

1. **Generalizing Dijkstra monads**, i.e., a program verification technique for arbitrary monadic effects;

2. **Relational reasoning in F⋆**: devising scalable verification techniques for properties of multiple program executions (e.g., confidentiality, noninterference) or of multiple programs (e.g., program equivalence).

3. **Making F⋆'s effect system more flexible**, by supporting tractable forms of effect polymorphism and allowing some of the effects of a computation to be hidden if they do not impact the observable behavior;

4. Working out more of the **F⋆ semantics and metatheory**

5. Solving the **engineering challenges** of building a usable verification system.

Tezos Interest   We believe that this proposal will have a positive impact on the Tezos ecosystem. There is significant interest at Nomadic Labs to use F⋆ for various program verification tasks. In particular, Tezos is already using and increasingly also contributing to the HACL⋆ cryptographic library that is developed by Prosecco and formally verified in F⋆. Moreover, Marco Stronati (from Nomadic Labs, previously in Prosecco) made recent experiments implementing a simple variant of the Tezos economic protocol in F⋆. Santiago Zanella (from Microsoft Research) has also recently developed an F⋆ interpreter and program logic for Michelson, the low-level contract language of Tezos, and there seems to be a general interest in the community in verifying smart contracts written in Michelson or various higher-level contract languages. Finally, another good target for formal verification using F⋆ is a compiler between such a higher-level contract language and Michelson. We believe that all these Tezos projects could significantly benefit from having a faster evolution of F⋆, better semantic foundations, and a better engineered tool, which are the focus of the current proposal.

1   Generalizing Dijkstra Monads

One of the key distinguishing features of F⋆ compared to proof assistants like Coq is the first-class treatment of effects in a way that enables efficient automatic reasoning. Dijkstra monads [31, 32] are the mechanism by which F⋆ efficiently computes verification conditions, generically for any effect. While the verification condition generation algorithm is generic, for each effect one needs to define monadic operations used to combine weakest preconditions (return, bind, and any effect-specific actions). Our recent work [3] shows that for a certain class of effects it is possible to derive these operations automatically and prove their correctness generically, starting from a monadic definition of the effect itself. The class of effects for which we can currently derive Dijksta monads is, however, restrictive: the monadic definition of the effect has to fit a small syntactic subset of F⋆ in order for our syntactic translation and generic logical relation proof to apply. While the class of supported effects includes state and exceptions, it excludes other important effects such as non-termination, IO, nondeterminism, and probabilities.

We are currently working on lifting these limitations by developing a general semantic framework for verifying programs with arbitrary monadic side-effects using Dijkstra monads. The key insight is that any monad morphism between a computational monad and a specification monad gives rise to a Dijkstra monad, which provides great flexibility for obtain-

ing Dijkstra monads tailored to the verification task at hand. We believe that a large variety of specification monads can be obtained by applying monad transformers to various base specification monads, including predicate transformers and Hoare-style pre- and postconditions. For simplifying the task of defining correct monad transformers, we are working on a language inspired by Moggi's monadic metalanguage that is parameterized by a dependent type theory. And we are also trying to develop a notion of Plotkin and Power's algebraic operations for Dijkstra monads, together with a corresponding notion of effect handlers. Finally, we are implementing our framework in both F* and Coq, and illustrate that it supports a wide variety of verification styles for effects such as partiality, exceptions, nondeterminism, state, and input-output; and we hope to also extend this to probabilities.

This generic framework will bring important improvements to F*: (1) it will remove the previous restrictions on the supported side-effects to input-output, nondeterminism, and probabilities; (2) it will give us much more flexibility in choosing the specification and verification style best suited for the verification task at hand; (3) it will provide a principled way to support algebraic operations and effect handlers in F*. Finally, we believe we will also be able to extend Dijkstra monads and our semantic framework to relational reasoning, in order to obtain principled semi-automated verification techniques for properties of multiple program executions (e.g., noninterference) or of multiple programs (e.g., program equivalence).

2   Relational reasoning in F*

By default, F* reasons about the execution of each effectful computation intrinsically, when the computation is defined, by inferring a (unary) weakest precondition (WP) for it. In recent work [17], we showed that by exposing the representation of the effect via monadic reification we can also reason about effectful computations extrinsically, after the fact. This allows proving *relational properties*, describing multiple executions of one or more programs. We evaluated this idea by encoding a variety of relational program analyses, including information flow control, program equivalence and refinement at higher order, correctness of program optimizations and game-based cryptographic security.

While monadic reification worked reasonably well on simple examples, two serious challenges remain: (1) monadic reification seems very difficult to soundly combine with reasoning about monotonic state [2], which at the moment crucially relies on the effect being treated abstractly, and the programmer not having direct access to the heap in user-programs; and (2) monadic reification seems hard to scale to large verification efforts such as miTLS [9, 10], which currently relies on meta-level arguments involving parametricity that are not fully formal, but that are at least more modular. One way to approach problem (1) would be to figure out how to safely use reification to expose the update monad structure [4] underpinning monotonic state. The challenge here is that the most natural approach inspired by hybrid modal logics [5, 29] (in which one also "reifies" the modal logic specifications of monotonic state) inevitably leads to reification becoming a whole program (and typing judgement) transformation, which of course does not scale well. Another way to approach this could be to keep the monadic represen-

tation abstract but do the relational reasoning not at the level of expressions but at the level of WPs; which can work if the WPs involved fully specify the result of the computation. Finally, for approaching the modularity issue (2) we could try to combine relational reasoning by reification or on WPs with internalized reasoning about parametricity [8, 19, 23].

3   More flexible effect system

Another interesting research direction is making the F* effect system more flexible. A first limitation we could try to lift is the lack of effect polymorphism, which means that higher-order functions such as `List.map` are duplicated many times, once for each effect their function argument might have. With the recent introduction in F* of support for type classes and for resolving implicit arguments using tactics, one can hope to also support the overloading (i.e., ad-hoc polymorphism) of functions such as `map` by parameterizing them over a Dijkstra monad and verifying their code in a generic fashion using tactics to apply the properties of Dijkstra monads.

A second limitation is that at the moment, effects in F* are syntactic: if any sub-computation triggers a certain effect then the whole computation is tainted with that effect. An ambitious project would be to relax this when the effect of a computation is unobservable to its context [11, 20, 24]. For example, consider computations that use state locally for memoization, or those that handle all exceptions that may be raised: it would be convenient to treat such computations as pure. Proving that we can soundly forget such non-observable effects proved to be hard already for state [6, 7, 34], many works targeting the case of hidden state [26, 30]. In the ideal case, we should be able to forget most effects, provided we prove in F* that these effects do not matter (which often requires relational reasoning, which is a separate topic below).

In the longer term it would be nice to also be able to prove termination extrinsically and then hide the divergence effect. One first step towards this could be to make the purity and divergence effects less primitive in F*, by encoding partiality as a monad supporting fixpoints [12, 15, 22] and on top of this defining Dijkstra monads for total and partial correctness.

4   F* semantics and metatheory

The semantics and metatheory of F* are interesting and still work in progress. While our previous work formally investigated several subsets of F* [2, 3, 16, 31], to be tractable these subsets simplified away several F* features that are in fact interesting. We would like to extend the theory of F* to cover such features, with a particular focus on solving challenging conceptual issues that could directly inform the further evolution and design of the language.

One important issue that is easy to explain is that in F* the type `list nat` is not a subtype of `list int`, even though `nat` is a subtype of `int`. We have the intuition that such subtyping on datatypes could be safely allowed if we additionally prevent pattern matching on datatype parameters (since allowing these parameters to be projected out immediately leads to unsoundness: `https://github.com/FStarLang/FStar/issues/65`). However, validating that this is indeed sound requires working out the metatheory of an extensional type theory with inductive types, refinement types, and subtyping.

One could potentially take inspiration in the new universe cumulative inductive types in Coq [33] and other recent work on subtyping for inductive types [1].

Another interesting issue to formally investigate is the treatment of F⋆'s propositional universe (prop) and its squashing of types [25]. While several encodings of prop were investigated, they all rely on axioms (e.g., for eliminating squash types) that need to be semantically justified to ensure the logical consistency of F⋆. Moreover, practical F⋆ developments also rely on further axioms such as (not entirely standard variants of) functional and propositional extensionality, whose soundness also needs to be semantically justified. One way to approach this would be to prove consistency taking inspiration in established semantic model constructions for existing type theories [18] such as Martin-Löf type theory [13, 14]. This could shed more light on the connection between F⋆ and established systems like Coq, Agda or the PRL family. Beyond prop and squashing, the F⋆-specific challenges include equality reflection, semantic termination, subtyping, etc.

5 Solving engineering challenges of building a usable verification system

On the engineering side, we will help turn F⋆ into an usable verification system. We will focusing in particular to the needs of libraries like HACL⋆, in particular by adding support for refinement proofs, which would greatly reduce code duplication. In addition, we will improve the support for metaprogramming and tactics, revamp the higher-order unification implementation, and add support for concurrency to F⋆. We will improve the F⋆ and Low⋆ libraries and documentation, build a comprehensive unit test suite, do performance profiling, monitor performance regression, etc. While these kind of tasks are maybe less interesting research-wise, they are nevertheless very important, and they can help make F⋆ more usable in an industrial context in general, and at Nomadic Labs and Tezos in particular.

References

[1] A. Abel, A. Vezzosi, and T. Winterhalter. Normalization by evaluation for sized dependent types. *PACMPL*, 1(ICFP):33:1–33:30, 2017.

[2] D. Ahman, C. Fournet, C. Hriţcu, K. Maillard, A. Rastogi, and N. Swamy. Recalling a witness: Foundations and applications of monotonic state. *PACMPL*, 2(POPL):65:1–65:30, 2018.

[3] D. Ahman, C. Hriţcu, K. Maillard, G. Martínez, G. Plotkin, J. Protzenko, A. Rastogi, and N. Swamy. Dijkstra monads for free. *POPL*. 2017.

[4] D. Ahman and T. Uustalu. Update monads: Cointerpreting directed containers. *TYPES*, 2013.

[5] C. Areces and B. ten Cate. *Handbook of Modal Logic*, chapter Hybrid Logics, pages 821–868. Elsevier, 2007.

[6] N. Benton, M. Hofmann, and V. Nigam. Proof-relevant logical relations for name generation. *TLCA*. 2013.

[7] N. Benton, M. Hofmann, and V. Nigam. Abstract effects and proof-relevant logical relations. *POPL*. 2014.

[8] J. Bernardy, P. Jansson, and R. Paterson. Parametricity and dependent types. In P. Hudak and S. Weirich, editors, *Proceeding of the 15th ACM SIGPLAN international conference on Functional programming, ICFP 2010, Baltimore, Maryland, USA, September 27-29, 2010*. 2010.

[9] K. Bhargavan, B. Bond, A. Delignat-Lavaud, C. Fournet, C. Hawblitzel, C. Hriţcu, S. Ishtiaq, M. Kohlweiss, R. Leino, J. Lorch, K. Maillard, J. Pang, B. Parno, J. Protzenko, T. Ramananandro, A. Rane, A. Rastogi, N. Swamy, L. Thompson, P. Wang, S. Zanella-Béguelin, and J.-K. Zinzindohoué. Everest: Towards a verified, drop-in replacement of HTTPS. *SNAPL*, 2017.

[10] K. Bhargavan, A. Delignat-Lavaud, C. Fournet, M. Kohlweiss, J. Pan, J. Protzenko, A. Rastogi, N. Swamy, S. Zanella Béguelin, and J. K. Zinzindohoue. Implementing and proving the TLS 1.3 record layer. *IEEE Security & Privacy*, 2017.

[11] L. Birkedal, G. Jaber, F. Sieczkowski, and J. Thamsborg. A kripke logical relation for effect-based program transformations. *Inf. Comput.*, 249:160–189, 2016.

[12] A. Bove and V. Capretta. Modelling general recursion in type theory. *Mathematical Structures in Computer Science*, 15(4):671–708, 2005.

[13] S. Castellan. Dependent type theory as the initial category with families. Internship Report, 2014.

[14] P. Clairambault and P. Dybjer. The biequivalence of locally cartesian closed categories and Martin-Löf type theories. *Mathematical Structures in Computer Science*, 24(6), 2014.

[15] F. Faissole and B. Spitters. Synthetic topology in homotopy type theory for probabilistic programming (extended abstract). accepted for PPS17 and CoqPL17, 2017.

[16] S. Forest and C. Hriţcu. Micro-F* in F*. Inria internship report, 2015.

[17] N. Grimm, K. Maillard, C. Fournet, C. Hriţcu, M. Maffei, J. Protzenko, T. Ramananandro, A. Rastogi, N. Swamy, and S. Zanella-Béguelin. A monadic framework for relational verification: Applied to information security, program equivalence, and optimizations. *CPP*, 2018.

[18] B. Jacobs. *Categorical logic and type theory*, volume 141. Elsevier, 1999.

[19] N. R. Krishnaswami and D. Dreyer. Internalizing relational parametricity in the extensional calculus of constructions. In S. R. D. Rocca, editor, *Computer Science Logic 2013 (CSL 2013), CSL 2013, September 2-5, 2013, Torino, Italy*. 2013.

[20] J. M. Lucassen and D. K. Gifford. Polymorphic effect systems. In J. Ferrante and P. Mager, editors, *Conference Record of the Fifteenth Annual ACM Symposium on Principles of Programming Languages, San Diego, California, USA, January 10-13, 1988*. 1988.

[21] G. Martínez, D. Ahman, V. Dumitrescu, N. Giannarakis, C. Hawblitzel, C. Hriţcu, M. Narasimhamurthy, Z. Paraskevopoulou, C. Pit-Claudel, J. Protzenko, T. Ramananandro, A. Rastogi, and N. Swamy. Meta-F*: Proof automation with SMT, tactics, and metaprograms. *ESOP*. 2019.

[22] C. McBride. Turing-completeness totally free. *MPC*. 2015.

[23] G. Moulin. *Internalizing Parametricity*. PhD thesis, Chalmers University of Technology, Gothenburg, Sweden, 2016.

[24] D. A. Naumann. Observational purity and encapsulation. *FASE*, 2005.

[25] A. Nogin. Quotient types: A modular approach. *TPHOLs*. 2002.

[26] F. Pottier. Hiding local state in direct style: A higher-order anti-frame rule. *LICS*. 2008.

[27] J. Protzenko, B. Beurdouche, D. Merigoux, and K. Bhargavan. Formally verified cryptographic web applications in WebAssembly. In *2019 IEEE Symposium on Security and Privacy (SP)*, 2019.

[28] J. Protzenko, J.-K. Zinzindohoué, A. Rastogi, T. Ramananan-dro, P. Wang, S. Zanella-Béguelin, A. Delignat-Lavaud, C. Hriţcu, K. Bhargavan, C. Fournet, and N. Swamy. Verified low-level programming embedded in F*. *PACMPL*, 1(ICFP):17:1–17:29, 2017.

[29] J. Reed. *A Hybrid Logical Framework*. PhD thesis, Carnegie Mellon University, 2009.

[30] J. Schwinghammer, L. Birkedal, F. Pottier, B. Reus, K. Støvring, and H. Yang. A step-indexed kripke model of hidden state. *Mathematical Structures in Computer Science*, 23(1):1–54, 2013.

[31] N. Swamy, C. Hriţcu, C. Keller, A. Rastogi, A. Delignat-Lavaud, S. Forest, K. Bhargavan, C. Fournet, P.-Y. Strub, M. Kohlweiss, J.-K. Zinzindohoue, and S. Zanella-Béguelin. Dependent types and multi-monadic effects in F*. *POPL*. 2016.

[32] N. Swamy, J. Weinberger, C. Schlesinger, J. Chen, and B. Livshits. Verifying higher-order programs with the Dijkstra monad. *PLDI*. 2013.

[33] A. Timany and M. Sozeau. Cumulative inductive types in Coq. In H. Kirchner, editor, *3rd International Conference on Formal Structures for Computation and Deduction, FSCD 2018, July 9-12, 2018, Oxford, UK*. 2018.

[34] A. Timany, L. Stefanesco, M. Krogh-Jespersen, and L. Birkedal. A logical relation for monadic encapsulation of state: proving contextual equivalences in the presence of runST. *PACMPL*, 2(POPL):64:1–64:28, 2018.

[35] J.-K. Zinzindohoué, K. Bhargavan, J. Protzenko, and B. Beurdouche. HACL*: A verified modern cryptographic library. *CCS*, 2017.