

SECOMP

Efficient Formally Secure Compilers to a Tagged Architecture

Cătălin Hrițcu

Inria Paris

(currently on leave at MSR Redmond)


SECOMP

- **grant recently funded by European Research Council (ERC)**
 - "most prestigious" individual research grants in Europe
- **5 year research project at Inria Paris (2017-2021)**
 - hiring: 3 PhD students, 2 PostDocs, 1 Starting Researcher
- **new people starting officially in January 2017:**
 - Marco Stronati (PostDoc, working on privacy at Cornell Tech NY)
 - Guglielmo Fachini (Research Engineer, "pre-PhD intern")
- **more collaborators and community building**
 - visits, sabbaticals, new secure compilation workshop, etc.
- **project builds mainly on Micro-Policies and Yannis' work**
 - Yannis left Inria end of September, very unfortunate for us



The problem: devastating low-level attacks

- **1. inherently insecure low-level languages**
 - **memory unsafe**: any buffer overflow can be catastrophic allowing remote attackers to gain complete control
- **2. unsafe interoperability with lower-level code**
 - even code written in **safer high-level languages** has to interoperate with **insecure low-level libraries**
 - **unsafe interoperability**: all high-level safety guarantees lost
- **Today's languages & compilers plagued by low-level attacks**
 - **hardware** provides no appropriate security mechanisms
 - fixing this purely in software would be way **too inefficient**



Key enabler: Micro-Policies



- software-defined, hardware-accelerated, tag-based monitoring
- micro-policies are cool!
 - **low level + fine grained**: unbounded per-word metadata, checked & propagated on each instruction
 - **flexible**: tags and monitor defined by software
 - **efficient**: software decisions hardware cached
 - **expressive**: complex policies for secure compilation
 - **secure** and **simple** enough to verify security in Coq
 - **real**: FPGA implementation on top of RISC-V



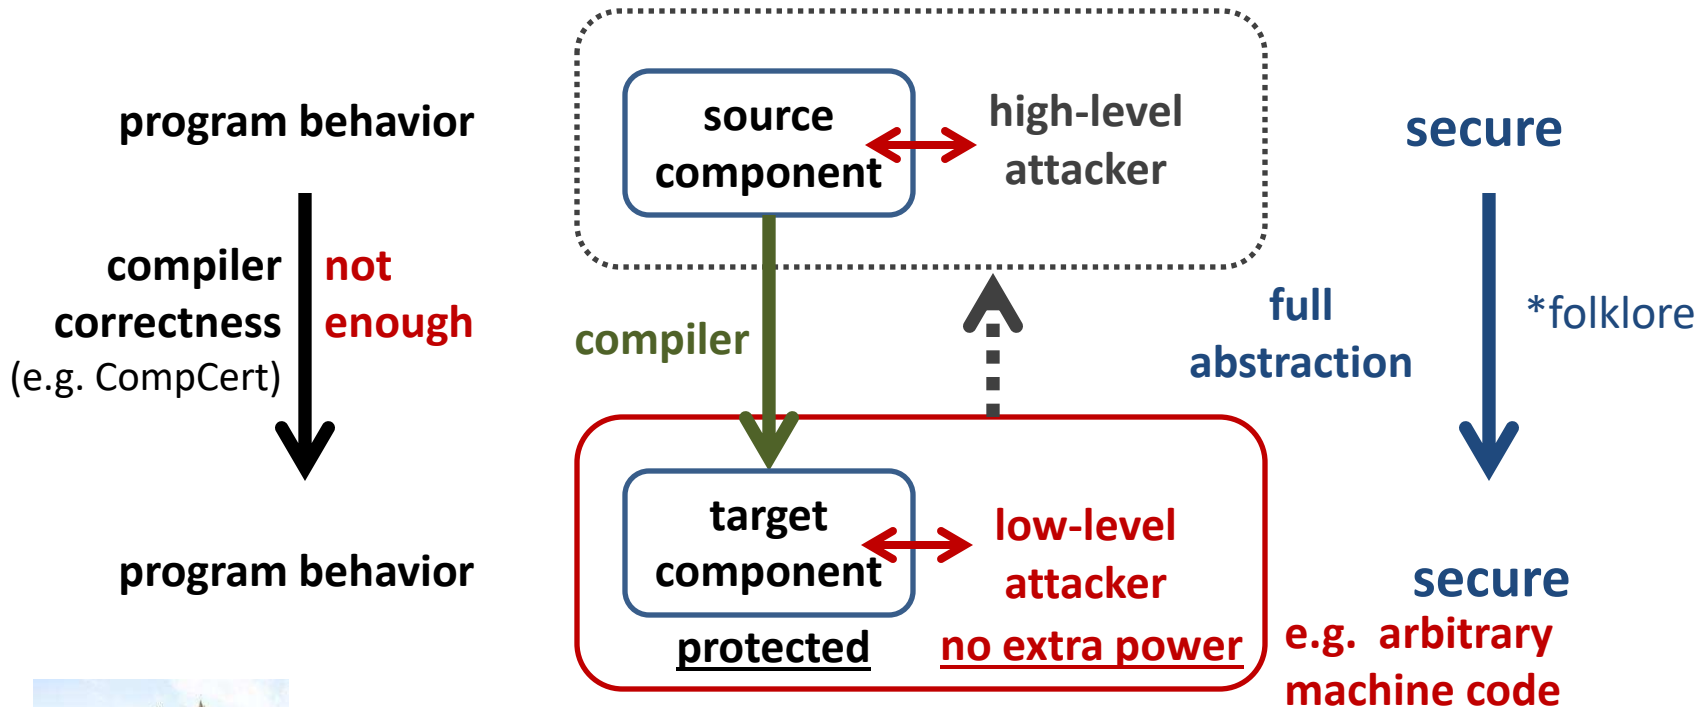
SECOMP grand challenge

Use micro-policies to build **the first efficient formally secure compilers** for **realistic programming languages**

1. **Provide secure semantics for low-level languages**
 - C with protected components and memory safety
2. **Enforce secure interoperability with lower-level code**
 - ASM, C, and F* [= ML + verification]

Formally verify: full abstraction

holy grail of secure compilation, enforcing abstractions all the way down



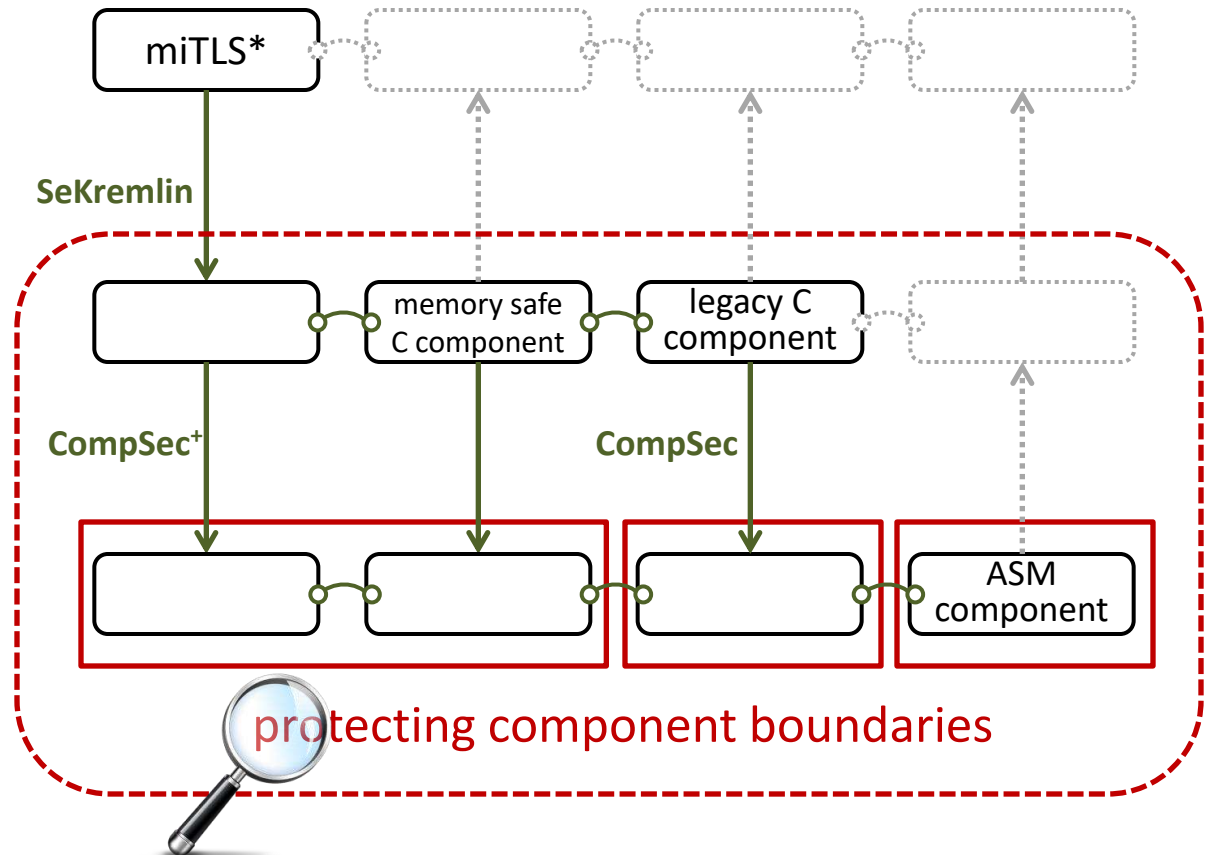
Benefit: sound security reasoning in the source language
forget about compiler chain (linker, loader, runtime system)
forget that libraries are written in a lower-level language

SECOMP: achieving full abstraction at scale

F* language
(ML + verification)

C language
+ memory safety
+ components

ASM language
(RISC-V + micro-policies)



Protecting component boundaries

- Break up software into **mutually distrustful components** running with **minimal privileges** & interacting only via **well-defined interfaces**
- **Limit the damage** of control hijacking attacks to just the C or ASM components where they occur
- Not a new idea, already deployed in practice:
 - process-level privilege separation
 - software-fault isolation
- Micro-policies can give us **better interaction model**
- We also aim to **show security formally**



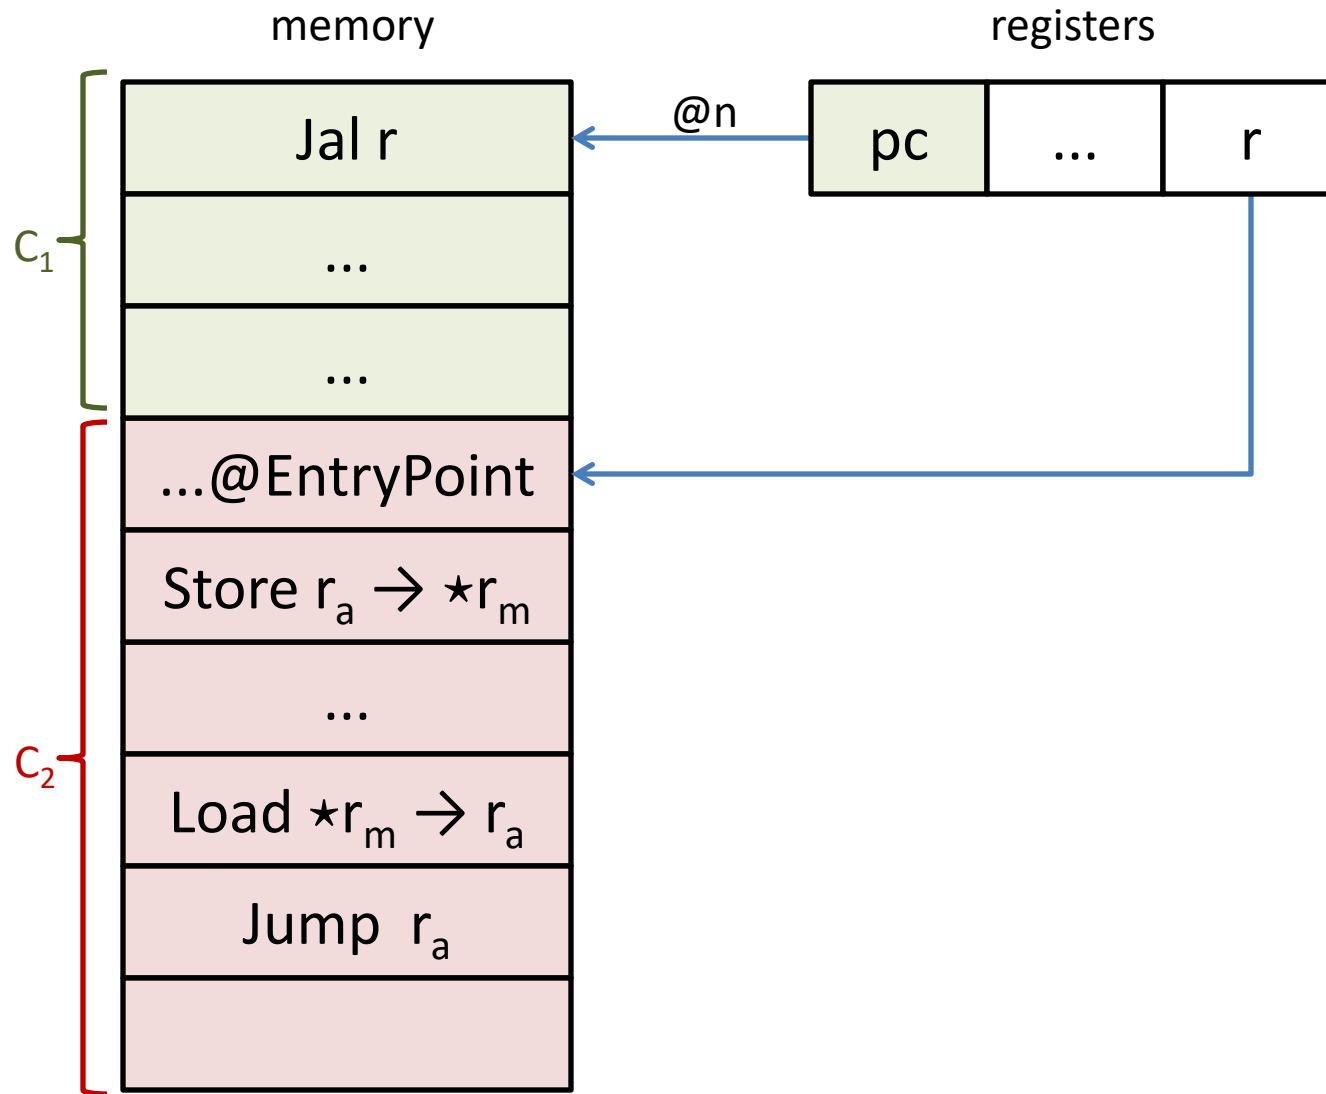
Towards compartmentalized C

- Want to **add components with typed interfaces to C**
- Compiler (e.g. CompCert), linker, loader propagate interface information to low-level memory tags
 - each component's memory tagged with unique color
 - procedure entry points tagged with procedure's type
- Micro-policy enforcing:
 - **component isolation**
 - **procedure call discipline** (entry points)
 - **stack discipline for returns** (linear return capabilities)
 - **type safety** on cross-component interaction

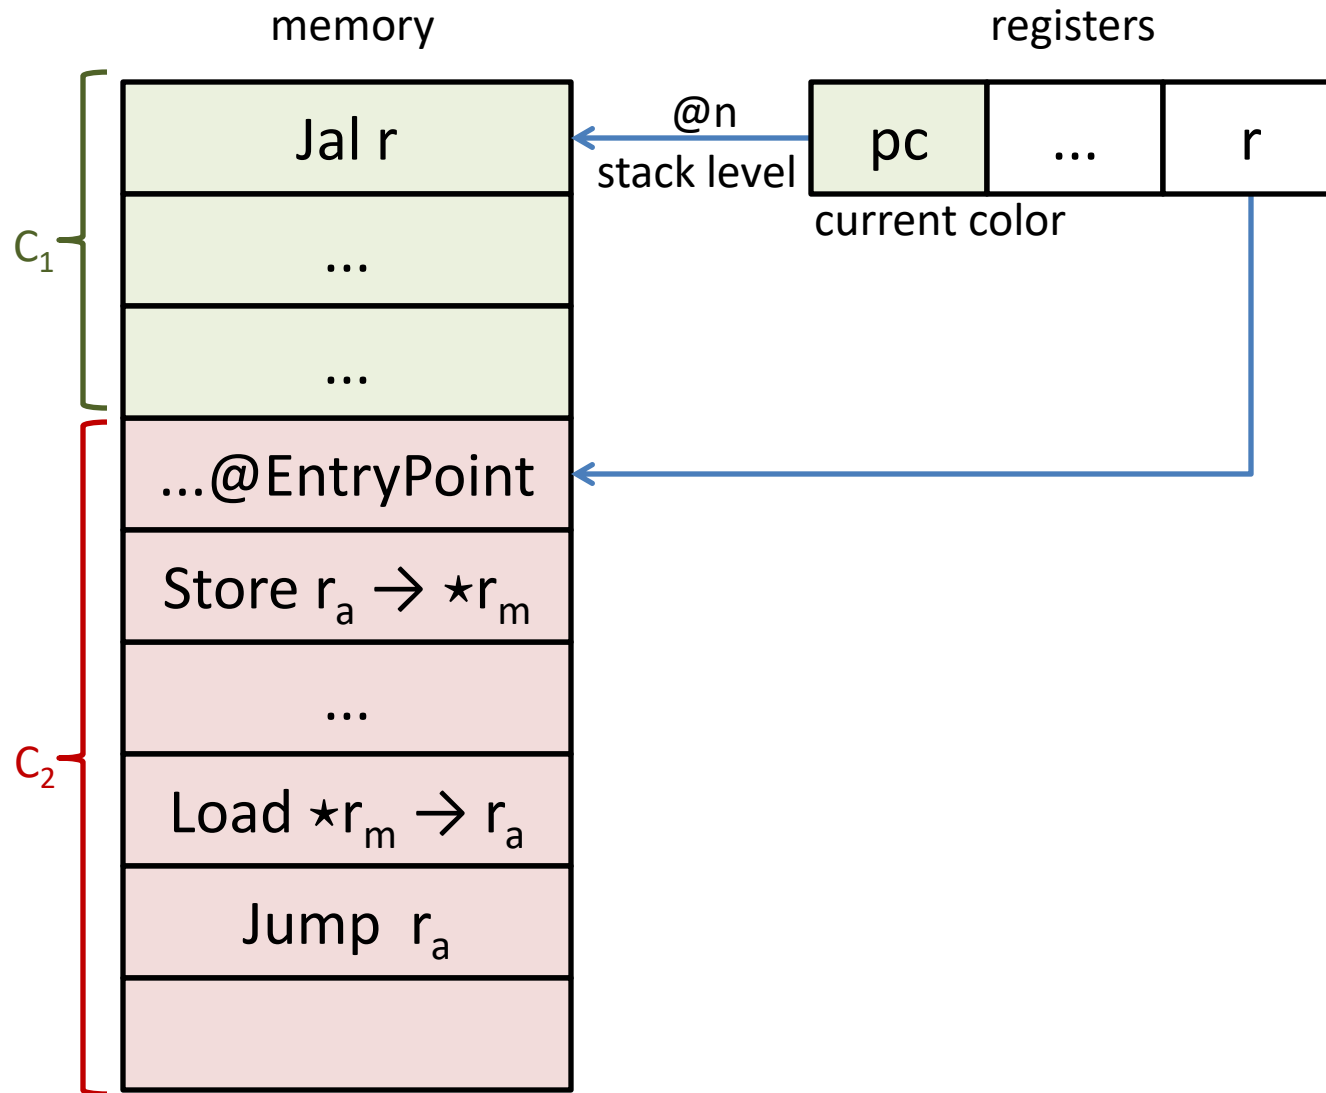


[Towards a Fully Abstract Compiler Using Micro-Policies, Yannis et al, TR 2015]

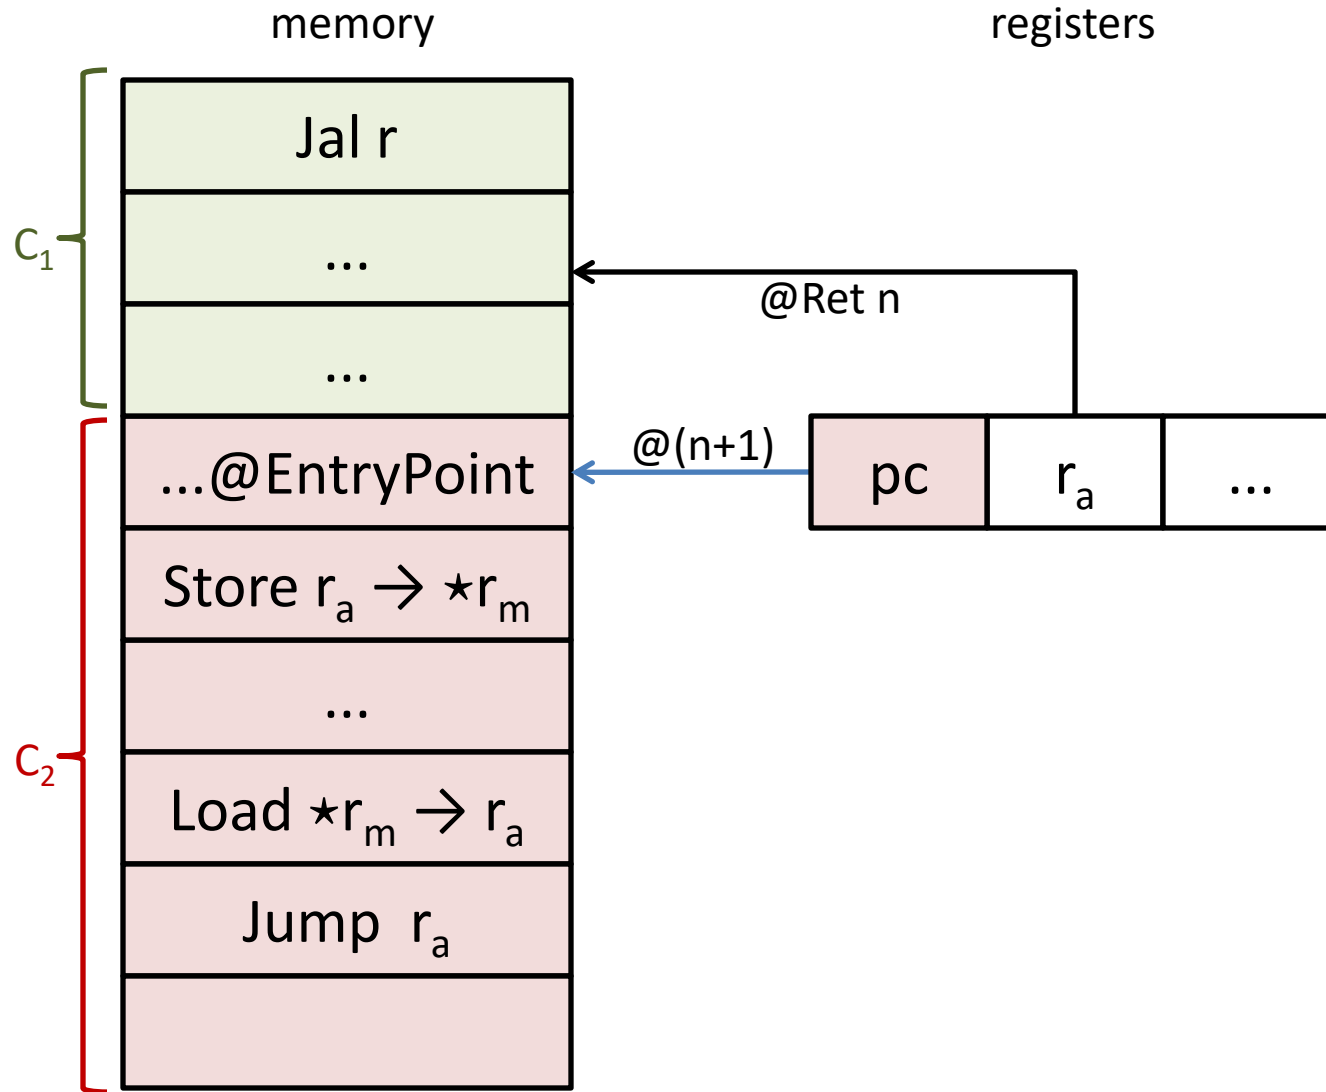
Compartmentalization micro-policy



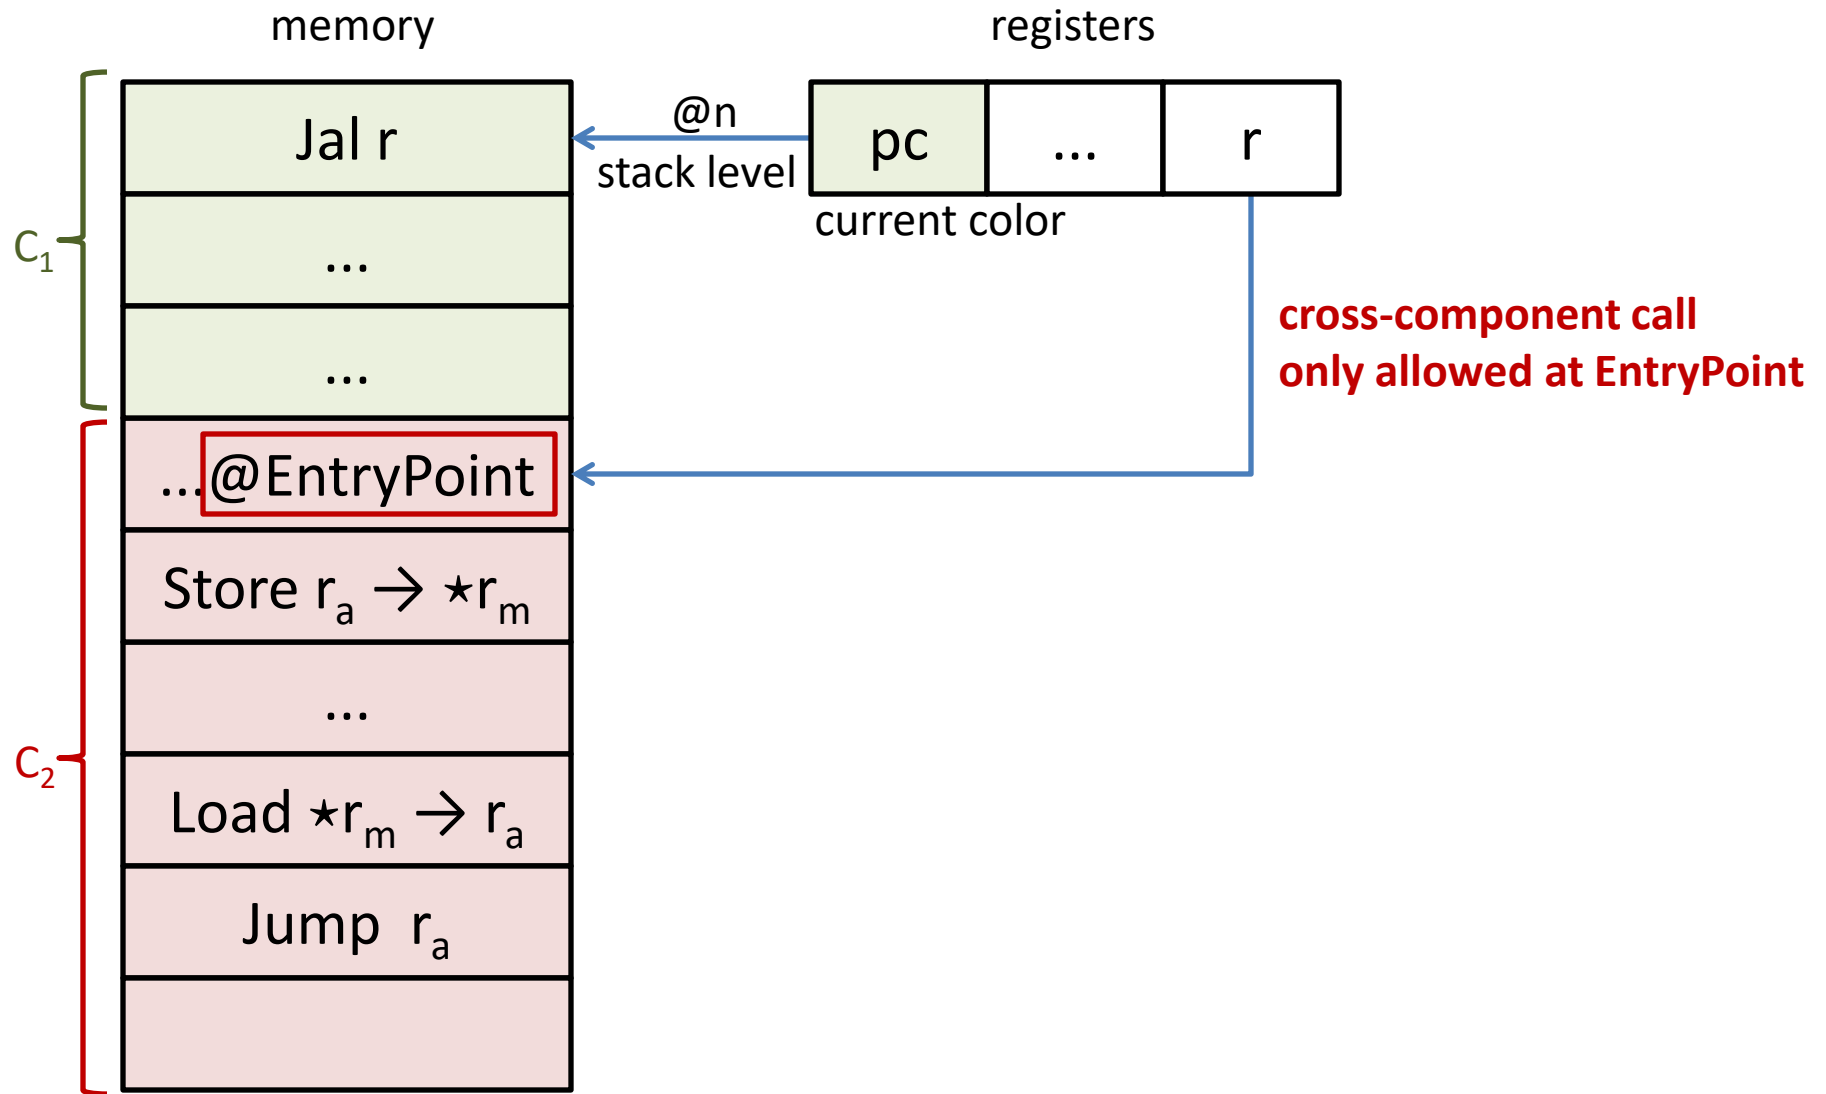
Compartmentalization micro-policy



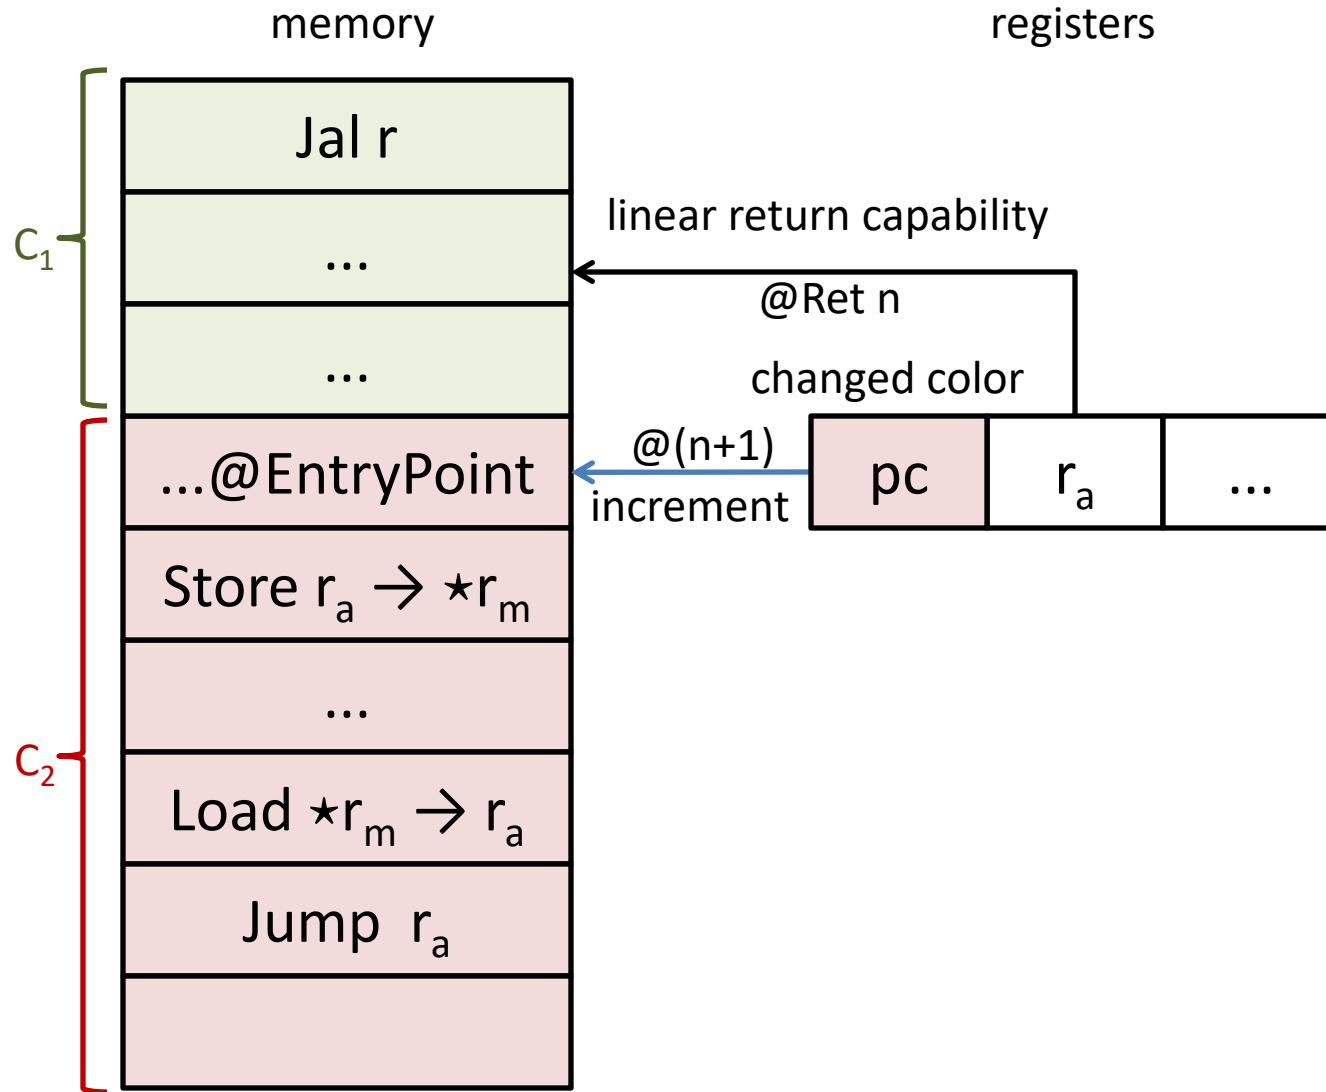
Compartmentalization micro-policy



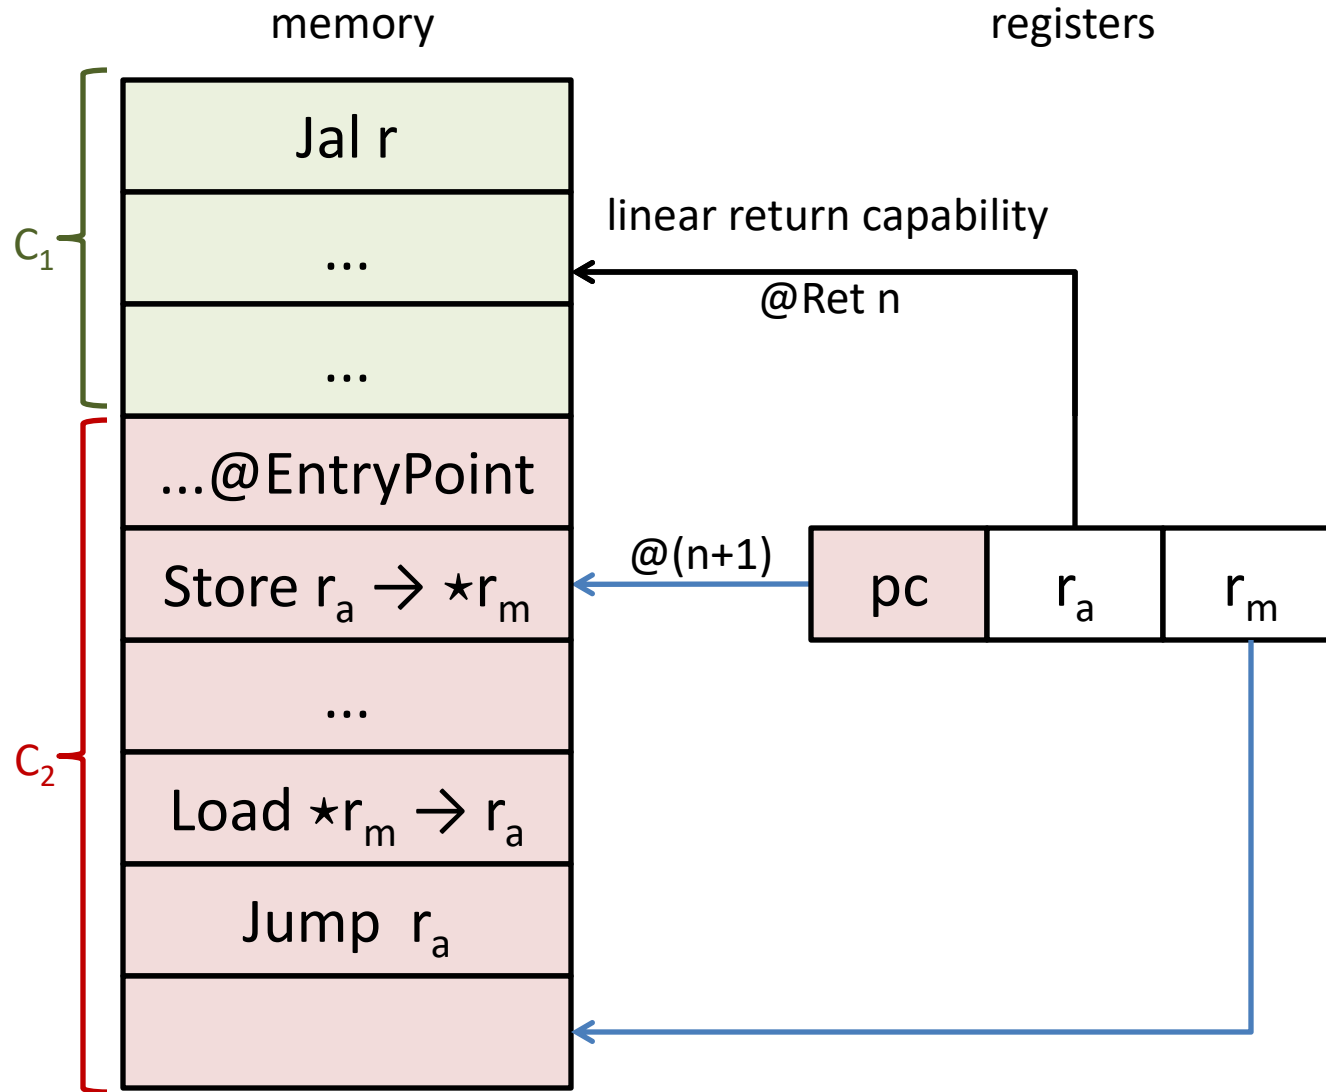
Compartmentalization micro-policy



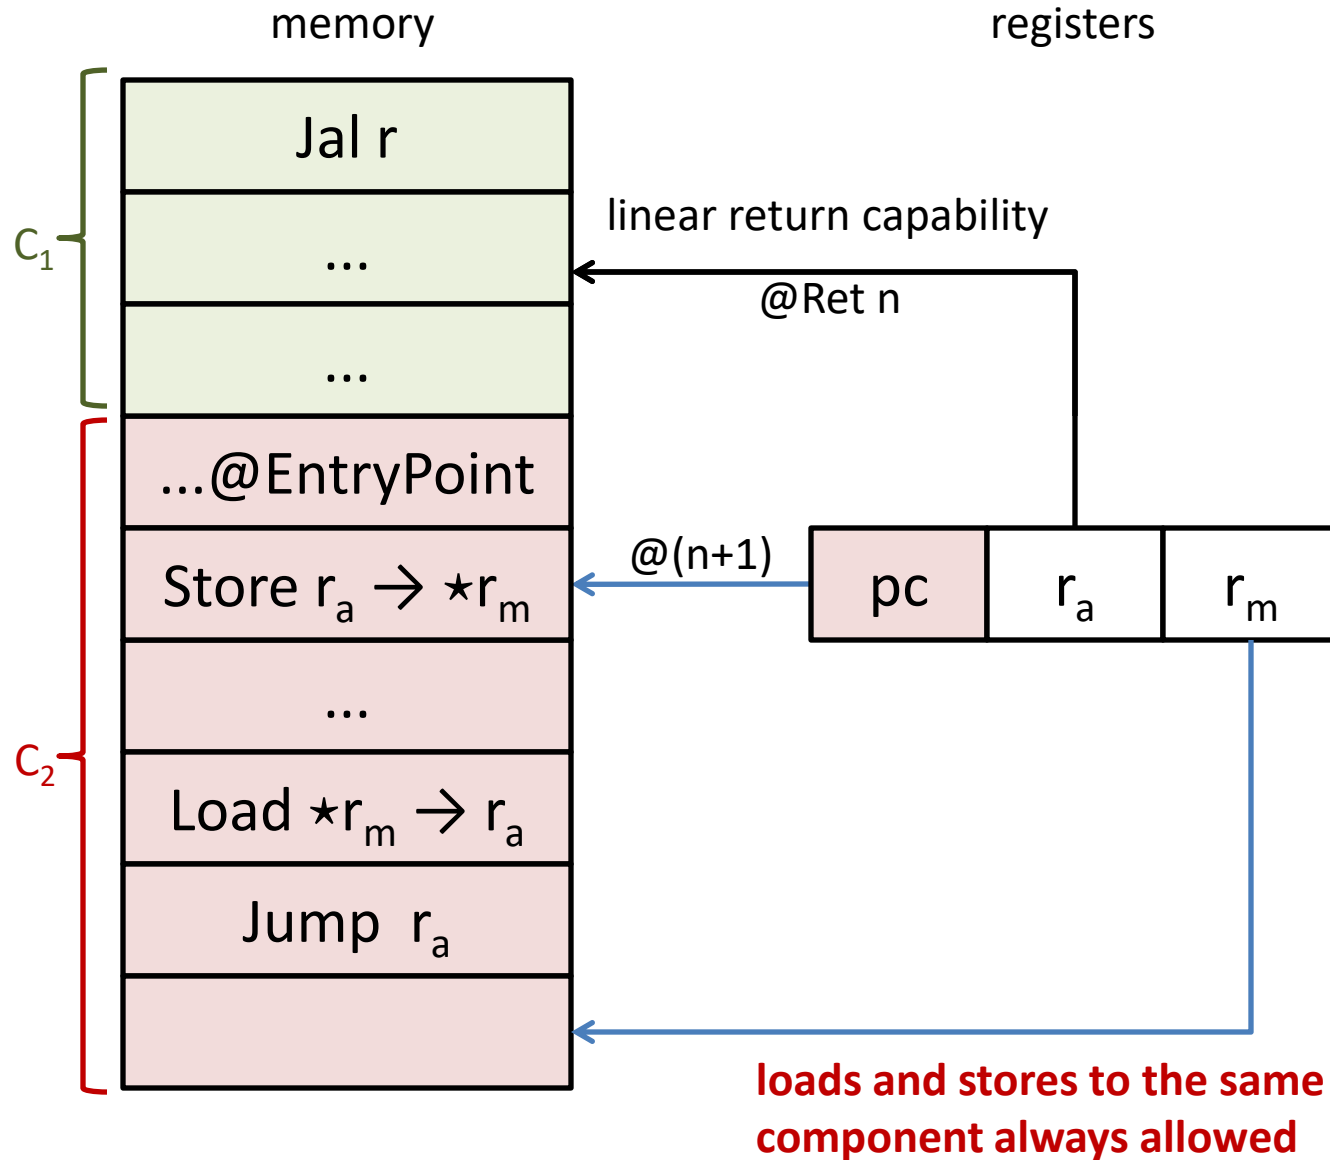
Compartmentalization micro-policy



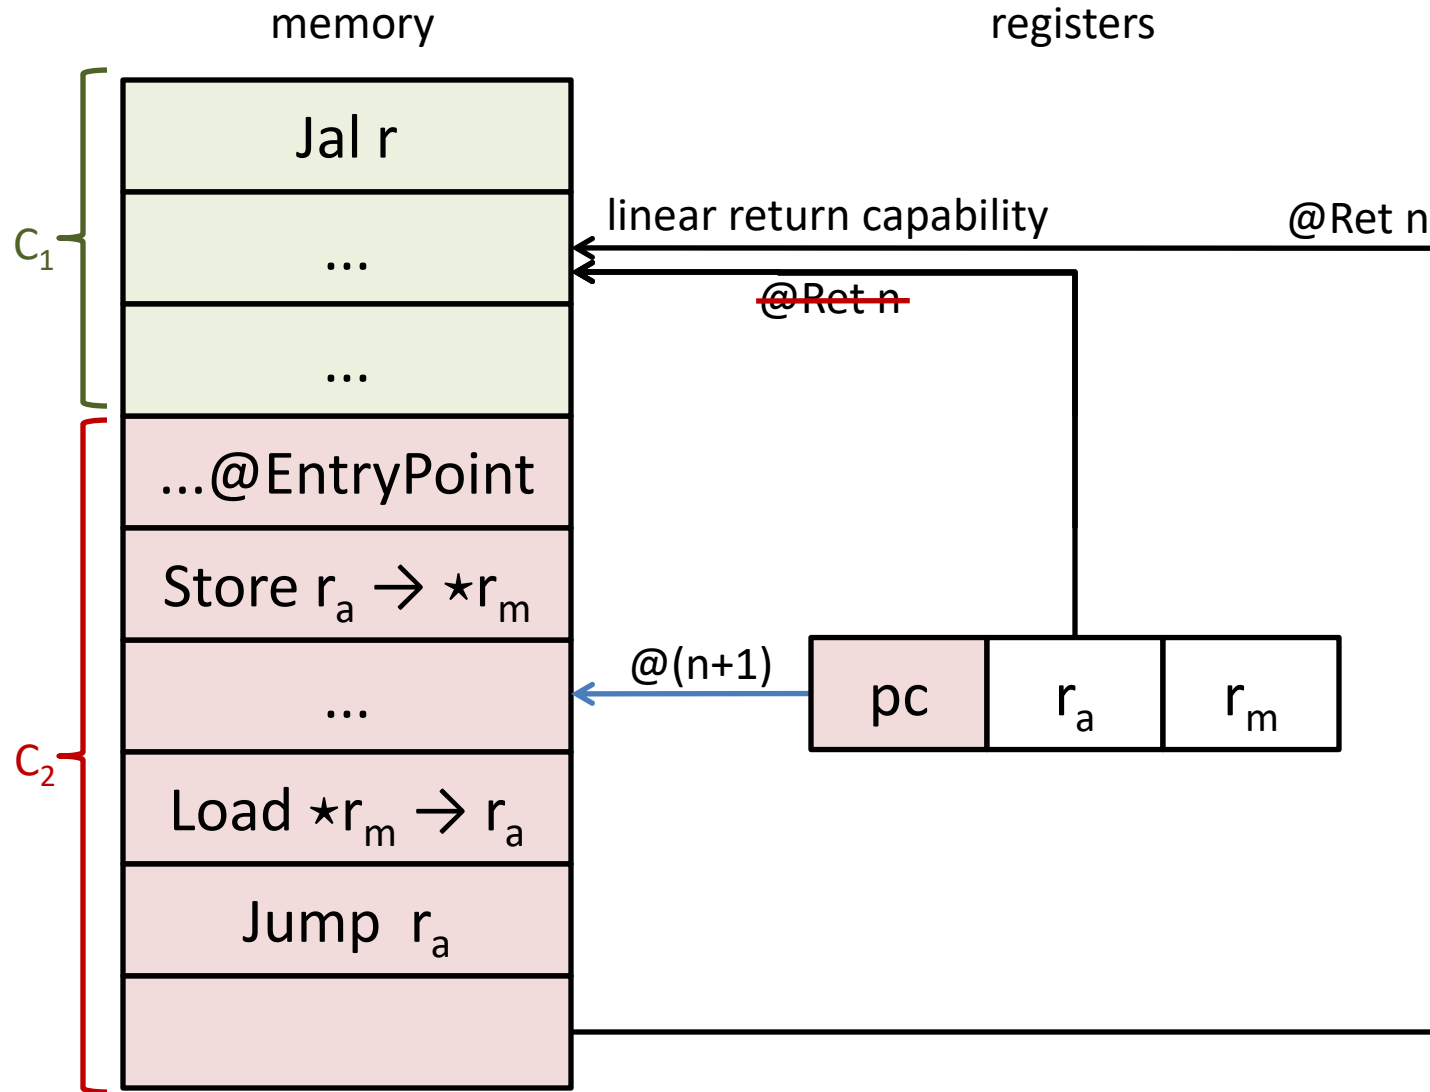
Compartmentalization micro-policy



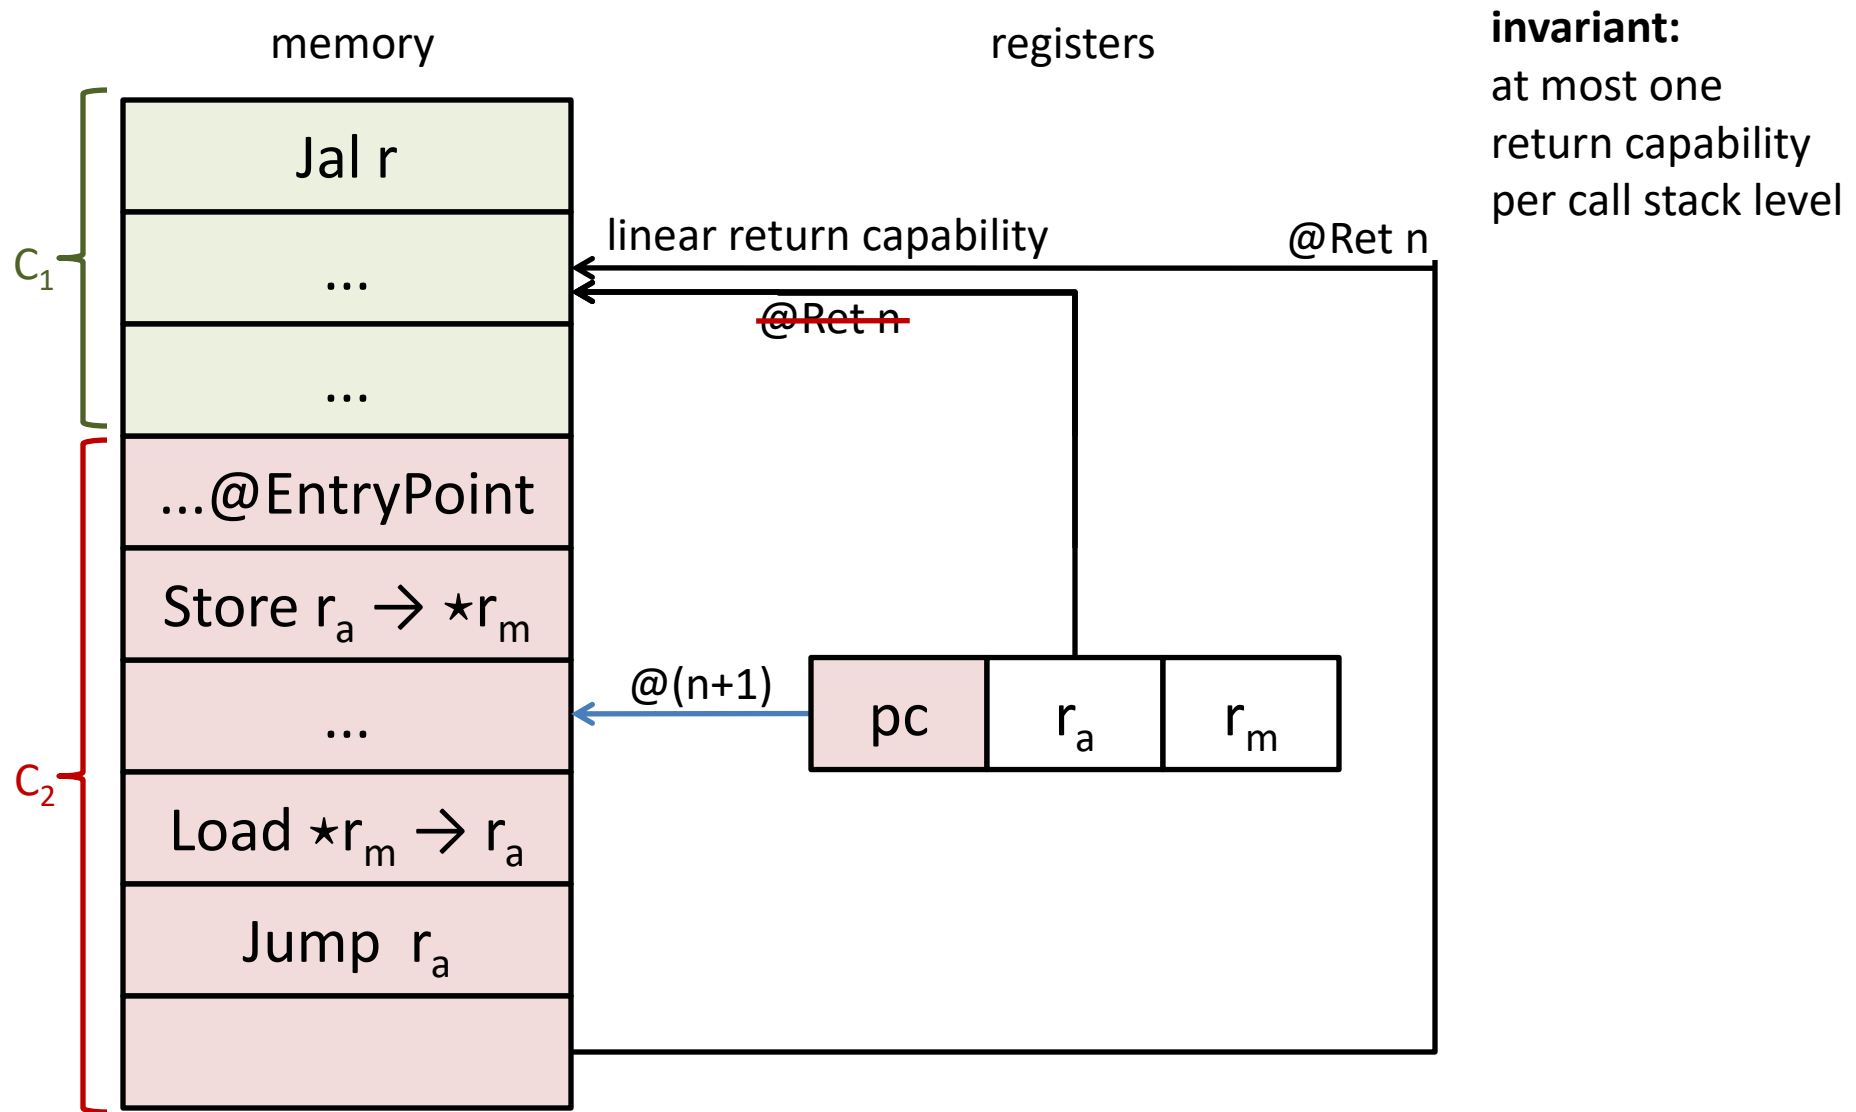
Compartmentalization micro-policy



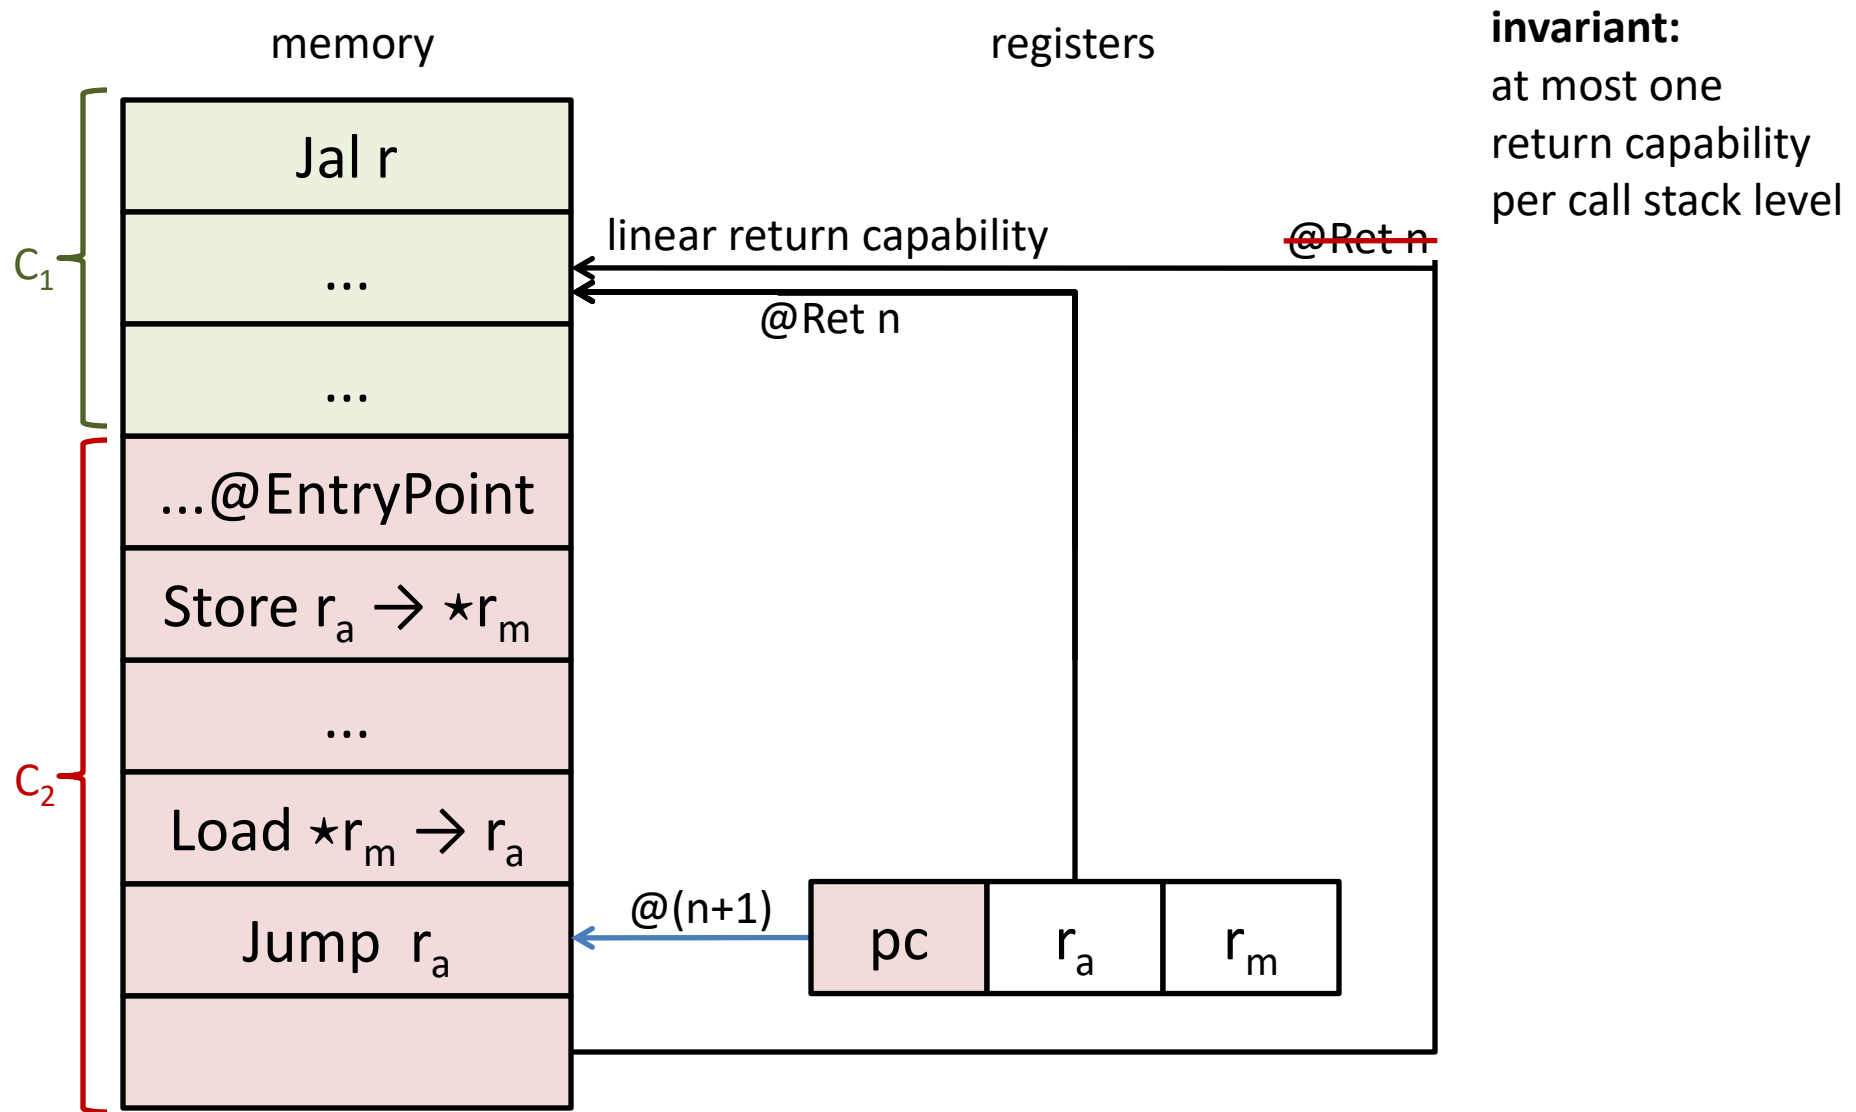
Compartmentalization micro-policy



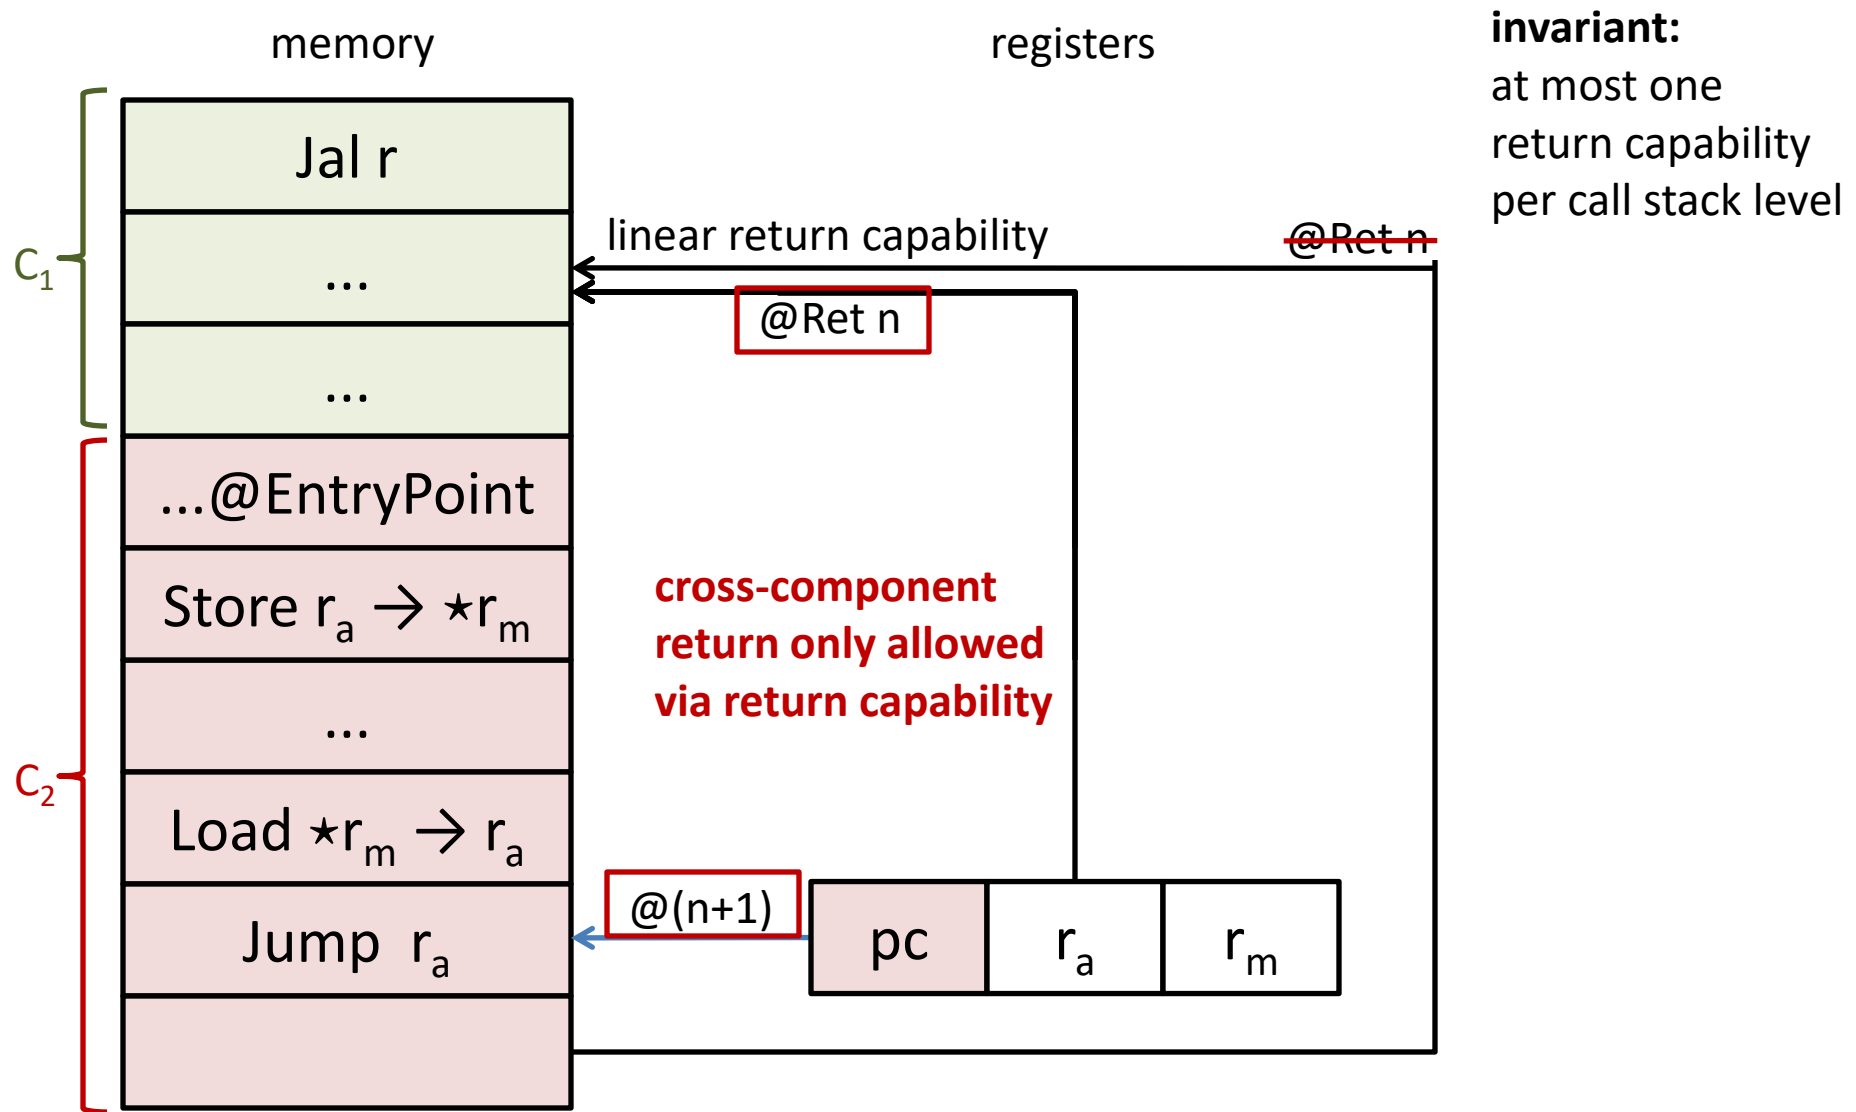
Compartmentalization micro-policy



Compartmentalization micro-policy



Compartmentalization micro-policy



Compartmentalization micro-policy?

	Yannis et al, TR 2015 and later	Antal et al, Oakland 2014 (SFI inspired)
abstraction level (source)	core C (and core Java)	machine code
compilation target	simple RISC machine code	
compartment lifetime	static	dynamic compartment creation
target property / attacker model	full abstraction variant [CSF 2016]	correct isolation (could be extended though)
mutual distrust justified by	unsafe source (C) + linking with unsafe / malicious machine code	interacting with unsafe / malicious machine code
enforced interaction model	valid calls and returns (μP); register cleaning and restoring; typed arguments and results (μP)	cross-compartment jumps only to designated entry points
memory protection	no cross-compartments writes or reads	sets of allowed cross- compartments reads and writes

Open problems on compartmentalization (1)

- **Dealing with more of C ... towards setjmp/longjmp ☺**
 - Yannis had **compartment-local stacks**
 - vs C: shared stack for all compartments, trickier
- **More hardware support could help**
 - **shared stack (+memory safety) require** setting tags on large regions of memory; in software (slow) or hardware (discussed)
 - **linear return capabilities** require PUMP inputs also be outputs
 - **cleaning registers** assumes compiler introduced instructions (restore) or some hardware support (all registers both inputs and outputs seems crazy, right?)

Open problems on compartmentalization (2)

- **Passing pointers between compartments**
 - currently can only allow immutable capabilities
 - e.g. code pointers as call capabilities
 - e.g. read/write capabilities to individual memory cells
 - capability is lost if pointer is changed
 - combining compartmentalization with memory safety allows richer object capability model
- **Linear return capabilities not transparent**
 - use wrapper or static analysis to gain transparency


Micro-policies:

remaining fundamental challenges

- **Micro-policies for C**
 - needed for vertical compiler composition
 - rule-based DSL for monitoring C programs
 - will put micro-policies in the hands of programmers
- **Secure micro-policy composition**
 - micro-policies are **interferent** reference monitors
 - one micro-policy's behavior can break another's guarantees
 - e.g. composing anything with IFC can leak

SECOMP in a nutshell

- We need more **secure languages, compilers, hardware**
- Key enabler: **micro-policies** (software-hardware protection)
- Grand challenge: **the first efficient formally secure compilers** for **realistic programming languages** (C and F*)
- Answering challenging fundamental questions
 - attacker models, proof techniques
 - secure composition, micro-policies for C
- Achieving strong security properties like full abstraction
 - + testing and proving formally that this is the case
- Measuring & lowering the cost of secure compilation



BACKUP SLIDES

SECOMP focused on dynamic enforcement **but static analysis could help too**



- **Improving efficiency**
 - removing spurious checks
 - just that by using micro-policies our compilers add few explicit checks
 - e.g. turn off memory safety checking for a statically memory safe component that never sends or receives pointers
- **Improving transparency**
 - allowing more safe behaviors
 - e.g. we could statically detect which copy of the linear return capability the code will use to return (in this case static analysis untrusted)

Beyond full abstraction

- Is full abstraction always the right notion of secure compilation? The right attacker model?
- **Similar properties**
 - secure compartmentalizing compilation (SCC)
 - preservation of hyper-safety properties [Garg et al.]
- **Strictly weaker properties** (easier to enforce!):
 - robust compilation (integrity but no confidentiality)
- **Orthogonal properties**:
 - memory safety (enforcing CompCert memory model)

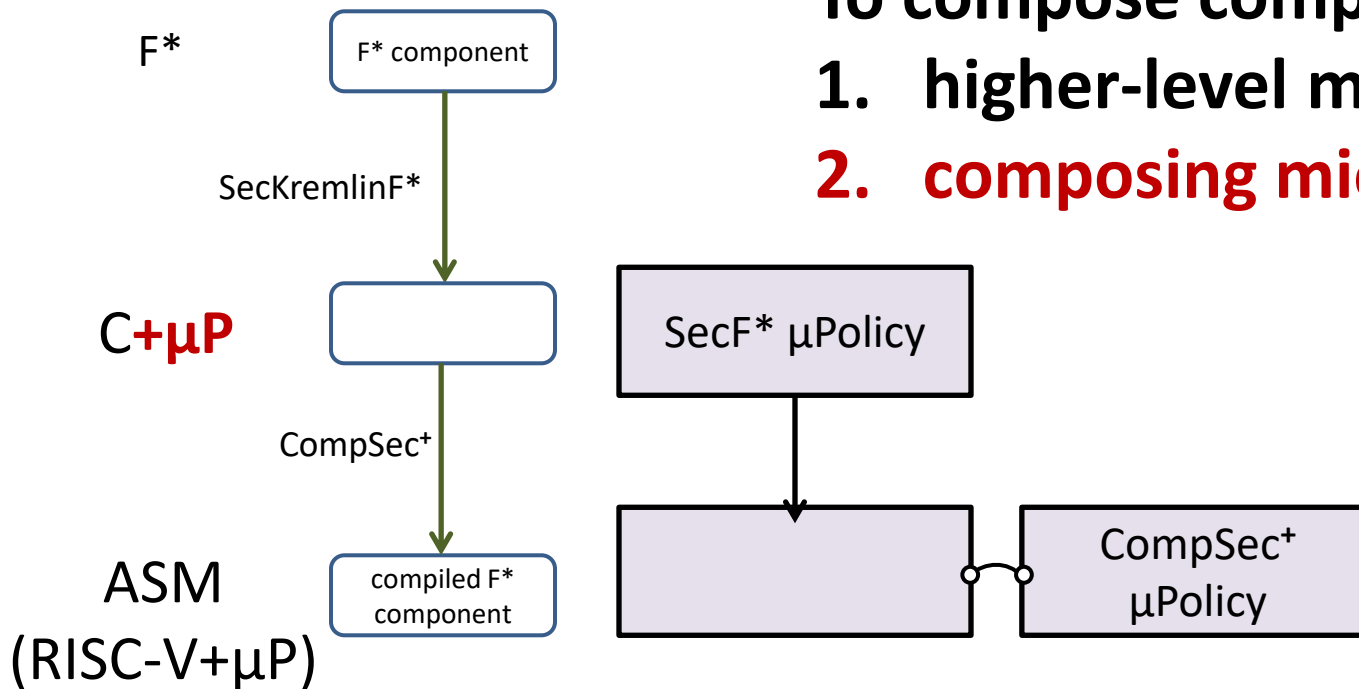
Collaborators & Community

- **Current collaborators from Micro-Policies project**
 - UPenn, MIT, Portland State, Draper Labs
- **Looking for additional collaborators**
 - Several other researchers working on **secure compilation**
 - Deepak Garg (MPI-SWS), Frank Piessens (KU Leuven), Amal Ahmed (Northeastern), Cedric Fournet & Nik Swamy (MSR)
- **Secure compilation meetings (very informal)**
 - 1st at INRIA Paris in August 2016
 - 2nd in Paris on 15 January 2017 before POPL at UPMC
 - **build larger research community, identify open problems, bring together communities** (hardware, systems, security, languages, verification, ...)

Composing compilers and higher-level micro-policies

To compose compilers need

1. higher-level micro-policies
2. **composing micro-policies**



User-specified higher-level policies

- By composing more micro-policies we can allow **user-specified micro-policies for C**
- Good news: **micro-policy composition is easy** since tags can be tuples
- But how do we ensure programmers won't break security?
- Bad news: **secure micro-policy composition is hard!**

