

Devastating low-level vulnerabilities

- Inherently insecure C/C++-like languages

- **memory (and type) unsafe:**
any buffer overflow is catastrophic
- **root cause**, but challenging to fix:
 - efficiency
 - precision
 - scalability
 - backwards compatibility
 - deployment



Practical mitigation: compartmentalization

- **Main idea:**

- break up security-critical C applications into **mutually distrustful components** running with **least privilege** & interacting via strictly enforced interfaces




- **Strong security guarantees & interesting attacker model**

- "a vulnerability in one component should not immediately destroy the security of the whole application"
- "components can be compromised by buffer overflows"
- "each component should be protected from all the others"

Goal 1: Formalize this

Goal 2: Build secure compilation chains

- **Add components to C**
 - interacting only via **strictly enforced interfaces**
- **Enforce "component C" abstractions:**
 - component separation, call-return discipline, ...
- **Secure compilation chain:**
 - compiler, linker, loader, runtime, system, hardware
- **Use efficient enforcement mechanisms:**
 - OS processes (all web browsers)
 - software fault isolation (SFI)
 - hardware enclaves (SGX)
 - WebAssembly (web browsers)
 - capability machines
 - tagged architectures
- **Practical need for this** (e.g. crypto library/protocol)



Source reasoning vs undefined behavior


- **Source reasoning**

= We want to reason formally about security
with respect to source language semantics


- **Undefined behavior**

= can't be expressed at all by source language semantics!


Dynamic compromise




→ \exists a **dynamic compromise scenario** explaining t in source language for instance leading to the following compromise sequence:



\wedge



\wedge



Trace is very helpful

- detect undefined behavior
- rewind execution




[When Good Components Go Bad - Fachini, Stronati, Hrițcu, et al]

Restricting undefined behavior

- **Mutually-distrustful components**
 - restrict **spatial** scope of undefined behavior
- **Dynamic compromise**
 - restrict **temporal** scope of undefined behavior
 - undefined behavior = **observable trace event**
 - **effects of undefined behavior**
 - shouldn't percolate before earlier observable events
 - careful with code motion, backwards static analysis, ...
 - CompCert **already offers** this saner model
 - GCC and LLVM **currently violate** this model

Now we know what these words mean!

(at least in the setting of compartmentalization for unsafe low-level languages)



Towards Secure Compilation Chain

(mostly)
Verified
(in Coq)



**Compartmentalized
unsafe source**



Buffers, procedures, components
interacting via **strictly enforced interfaces**

**Compartmentalized
abstract machine**



Simple RISC abstract machine with
build-in compartmentalization

**Micro-policy
machine**



Tag-based reference monitor enforcing:

- component separation
- procedure call and return discipline

(linear capabilities / linear entry points)

**Bare-bone
machine**

Inline reference monitor enforcing:

- component separation
- procedure call and return discipline

(program rewriting, shadow call stack)

software fault isolation

Systematically tested (with QuickChick)

