

Low-Level Software Security:

Attacks and Defenses;
Control Flow Integrity

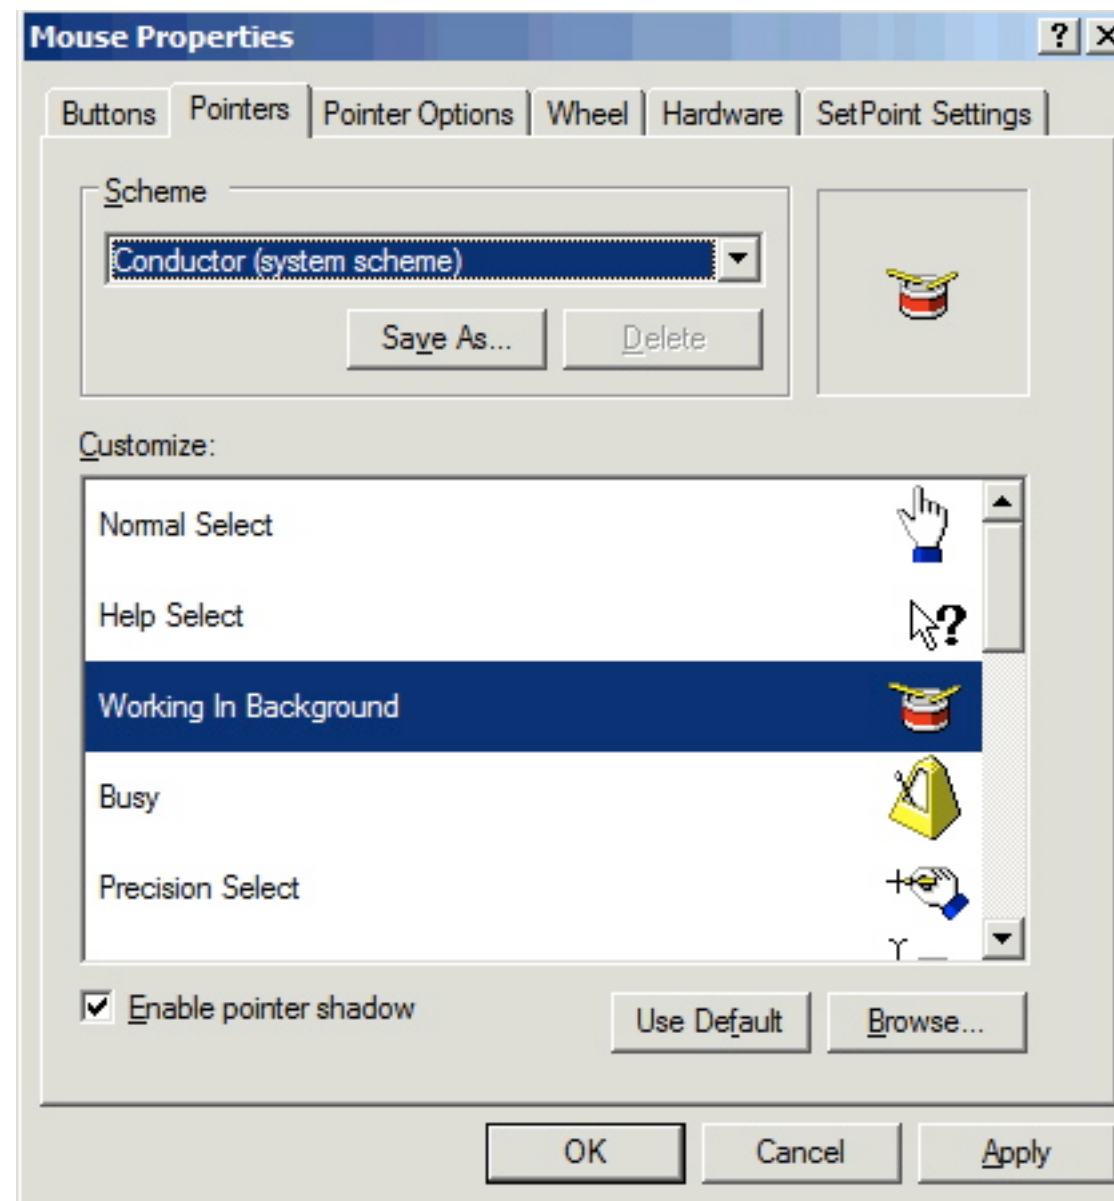
System Security Seminar
Presenter: Cătălin Hrițcu

References

- ▶ ***Low-Level Software Security: Attacks and Defenses***, by Úlfar Erlingsson (FOSAD 2007)
- ▶ ***Control-flow integrity***, by Martín Abadi, Mihai Budiu, Úlfar Erlingsson, Jay Ligatti (CCS 2005)
- ▶ Many thanks to Úlfar Erlingsson from Microsoft Research and Reykjavk University
 - ▶ Many of the slides are from his FOSAD 2007 talk
 - ▶ Opinions and mistakes are still mine

A real-world attack

- ▶ Microsoft Windows animated cursor buffer overflow vulnerability (March 30, 2007)



Rated as “Extremely critical”

[Home](#)[Corporate Website](#)[Jobs](#)[Mailing Lists](#)[RSS](#)[Blog](#)[Advert](#)

Solutions For

[Security Professionals](#)
[Security Vendors](#)

Free Solutions For


[Open Communities](#)
[Journalists & Media](#)

Online Services

[Secunia Blog](#)
[Online Software Inspector](#)
[Personal Software Inspector \(BETA\)](#)
[Network](#)

Microsoft Windows Animated Cursor Buffer Overflow Vulnerability

Secunia Advisory: SA24659
Release Date: 2007-03-30
Last Update: 2007-05-09

Critical:  [Extremely critical](#)
Impact: System access
Where: From remote
Solution Status: Vendor Patch

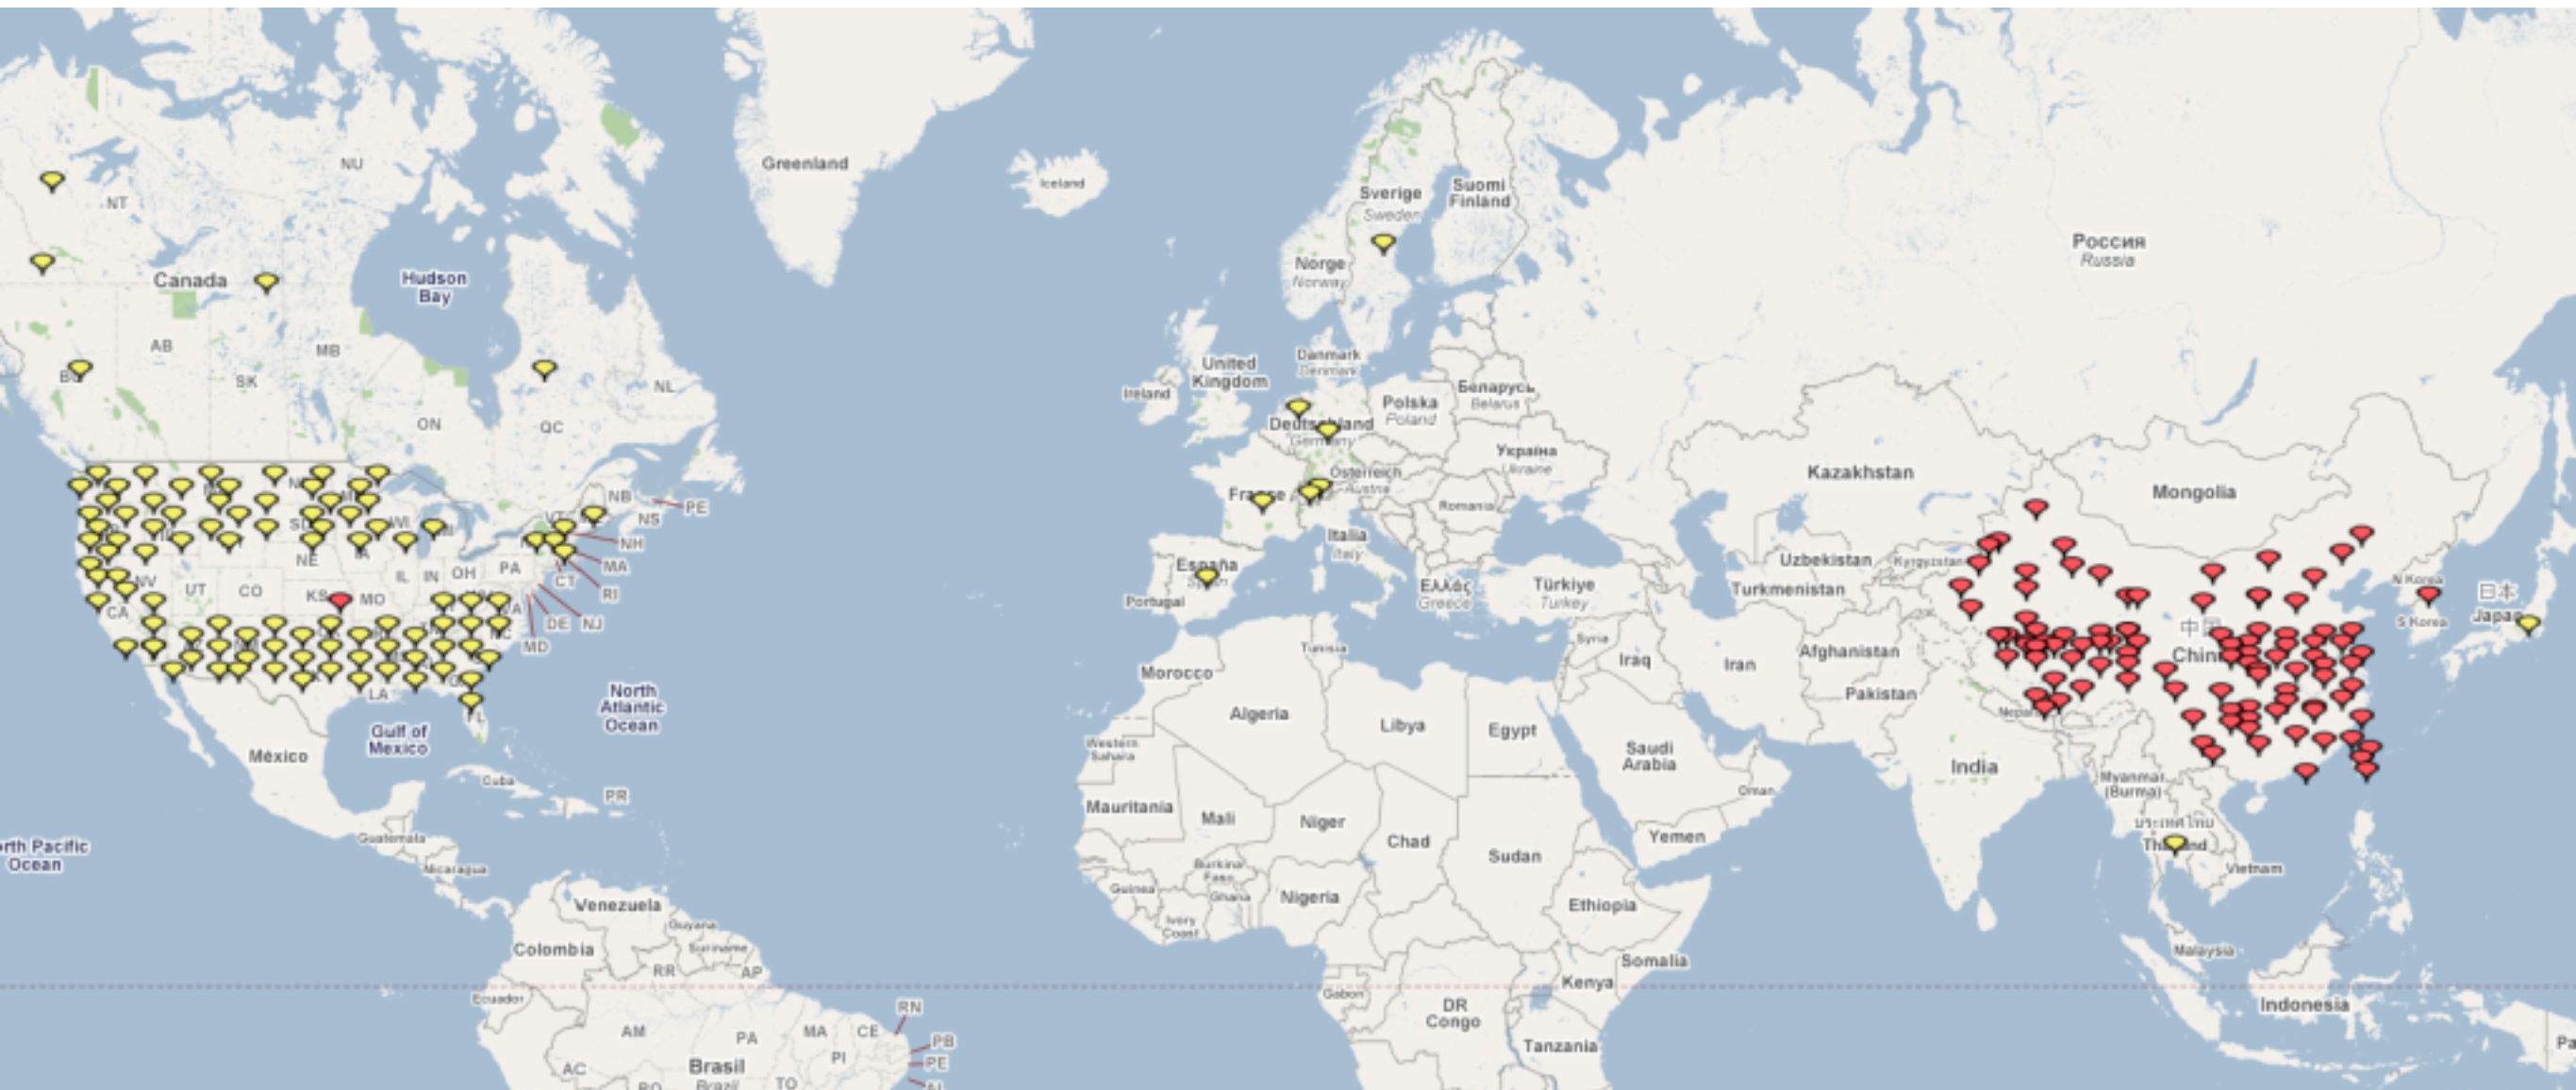
OS: [Microsoft Windows 2000 Advanced Server](#)
[Microsoft Windows 2000 Datacenter Server](#)
[Microsoft Windows 2000 Professional](#)
[Microsoft Windows 2000 Server](#)
[Microsoft Windows Server 2003 Datacenter Edition](#)
[Microsoft Windows Server 2003 Enterprise Edition](#)
[Microsoft Windows Server 2003 Standard Edition](#)
[Microsoft Windows Server 2003 Web Edition](#)
[Microsoft Windows Storage Server 2003](#)
[Microsoft Windows Vista](#)
[Microsoft Windows XP Home Edition](#)
[Microsoft Windows XP Professional](#)

Exploited immediately in the wild

► Attack vectors:

- ▶ HTML email / spam
- ▶ Browsers vulnerable - 2000+ different sites hosting exploit

From: Nude BritneySpeers.com [mailto:nudebritneyspeers.com]
Sent: Sunday, April 01, 2007 3:20 PM
To: [\[REDACTED\]](#)
Subject: Hot pictures of Britney Speers
Importance: High



Microsoft reacts fast

- Releases “out of band” patch (April 3)

Severity Ratings and Vulnerability Identifiers:


Vulnerability Identifiers	Impact of Vulnerability	Windows 2000 Service Pack 4	Windows XP Service Pack 2	Windows Server 2003, Windows Server 2003 Service Pack 1, and Windows Server 2003 Service Pack 2	Windows Vista
GDI Local Elevation of Privilege Vulnerability - CVE-2006-5758	Elevation of Privilege	Important	Important	Not Affected	Not Affected
WMF Denial of Service Vulnerability CVE-2007-1211	Denial of Service	Moderate	Moderate	Moderate	Not Affected
EMF Elevation of Privilege Vulnerability CVE-2007-1212	Elevation of Privilege	Important	Important	Important	Important
GDI Invalid Window Size Elevation of Privilege Vulnerability CVE-2006-5586	Elevation of Privilege	Important	Important	Not Affected	Not Affected
Windows Animated Cursor Remote Code Execution Vulnerability - CVE-2007-0038	Remote Code Execution	Critical	Critical	Critical	Critical
GDI Incorrect Parameter Local Elevation of Privilege Vulnerability - CVE-2007-1215	Elevation of Privilege	Important	Important	Important	Important
Font Rasterizer Vulnerability - CVE-2007-1213	Elevation of Privilege	Important	Not Affected	Not Affected	Not Affected
Aggregate Severity of All Vulnerabilities		Critical	Critical	Critical	Critical

- Unpatched systems are still being actively exploited

Low-level attacks

- ▶ Buffer overflows ~50% of all reported attacks
 - ▶ 1988: Robert Morris's Internet Worm
 - ▶ 2000: Code Red, SQL Slammer
 - ▶ 2007: Windows .ANI vulnerability
- ▶ Possible whenever compiling a high-level language into a low-level one without guaranteeing that the translation preserves the high-level abstractions

The legacy of C



- ▶ Millions of lines of security-critical code written in C
- ▶ C compilers do not guarantee any property of the translation
- ▶ Even if we ignore compilation, C is not type and memory safe
 - ▶ This often leads to exploitable vulnerabilities
- ▶ Safer low-level languages: Cyclone, CCured, SafeC, ...
 - ▶ Most of them are safe dialects of C
 - ▶ Existing C code can be migrated with some changes

Advertisement: Deputy

- ▶ Promising safe C compiler from Berkeley
 - ▶ Designed to work on existing real-world C code
 - ▶ Including the Linux kernel itself!
 - ▶ Allows for incremental transition
 - ▶ Memory layout of data structures preserved
 - ▶ Dependent types (additional annotations)
 - ▶ Hybrid type checking (static + dynamic)
 - ▶ Low performance impact
 - ▶ Average of 10-20% for CPU-intensive tests
 - ▶ Low annotation burden
 - ▶ For the Linux kernel less than 1% of lines annotated
- ▶ Might be the subject of a future talk in this seminar

High-level languages

- ▶ e.g. Java, C#, Python, O'Caml, Standard ML, ...
- ▶ Usually promise to be safer and more secure
- ▶ Their actual safety still depends on low-level details
 - ▶ Compiler and runtime-system need to be correct
 - ▶ Static checking is not enough
 - ▶ Most high-level properties need to be enforced at run-time
 - ▶ Safety vs. performance tradeoff
- ▶ Rewriting legacy code in a new language is in most cases out of the question
 - ▶ Even if the target language would be perfectly secure
 - ▶ And the run-time overhead negligible

Mitigation techniques

- ▶ The main subject of this talk
- ▶ Limited language-based defenses that
 - ▶ Work on legacy code
 - ▶ Are fully automatic
 - ▶ Operate at the lowest level (machine-code)
 - ▶ Involve no source-code changes (at most re-compilation)
 - ▶ Typically, runtime checks to guarantee high-level properties
 - ▶ Unobtrusive: close to **zero overhead** and zero false positives
 - ▶ Only prevent certain vulnerabilities / attacks
 - ▶ Often unclear what vulnerabilities are covered

Mitigations are a compromise

- ▶ Mitigations are limited, correct software is better
 - ▶ Wouldn't need any defenses if software was "correct"...
- ▶ So why not just fix all software?
 - ▶ Fixing software is difficult, costly, and error-prone
 - ▶ It is hard even to specify what "correct" should mean
 - ▶ Needs source, build environments, etc., and may interact badly with testing, debugging, deployment, and servicing
- ▶ Mitigations are not the optimal solution, but ...
 - ▶ They can rule out many practical attacks
 - ▶ Some of the ones we will see are deployed in practice (e.g. Windows XP SP2 and Vista)

Outline

- ▶ A simple buffer overflow example
- ▶ Two simple mitigation techniques
 - ▶ Stack canaries and cookies
 - ▶ Preventing data execution (NXD)
- ▶ A more complex jump-to-libc attack
- ▶ A more powerful mitigation technique
 - ▶ Control-flow integrity

A concrete stack overflow example

```
int is_file_foobar( char* one, char* two )
{
    // must have strlen(one) + strlen(two) < MAX_LEN
    char tmp[MAX_LEN];
    strcpy( tmp, one );
    strcat( tmp, two );
    return strcmp( tmp, "file://foobar" );
}
```

- ▶ Attack overflows a (fixed-size) array on the stack
- ▶ The function return address points to the attacker's code
- ▶ The best known low-level attack
 - ▶ Used by the Internet Worm in 1988 and commonplace since

A concrete stack overflow example

<u>address</u>	<u>content</u>	
0x0012ff5c	0x00353037	; argument two pointer
0x0012ff58	0x0035302f	; argument one pointer
0x0012ff54	0x00401263	; return address
0x0012ff50	0x0012ff7c	; saved base pointer
0x0012ff4c	0x00000072	; tmp continues 'r' '\0' '\0' '\0'
0x0012ff48	0x61626f6f	; tmp continues 'o' 'o' 'b' 'a'
0x0012ff44	0x662f2f3a	; tmp continues ':' '/' '/' 'f'
0x0012ff40	0x656c6966	; tmp array: 'f' 'i' 'l' 'e'

- ▶ The above stack snapshot is normal w/o overflow
- ▶ The arguments here are "file://" and "foobar"

A concrete stack overflow example

- ▶ A stack snapshot with a benign overflow


<u>address</u>	<u>content</u>	
0x0012ff5c	0x00353037	; argument two pointer
0x0012ff58	0x0035302f	; argument one pointer
0x0012ff54	0x00666473	; return address 's' 'd' 'f' '\0'
0x0012ff50	0x61666473	; saved base pointer 's' 'd' 'f' 'a'
0x0012ff4c	0x61666473	; tmp continues 's' 'd' 'f' 'a'
0x0012ff48	0x61666473	; tmp continues 's' 'd' 'f' 'a'
0x0012ff44	0x612f2f3a	; tmp continues ':' '/' '/' 'a'
0x0012ff40	0x656c6966	; tmp array: 'f' 'i' 'l' 'e'

- ▶ In the above, the stack has been corrupted
- ▶ The second (attacker-chosen) arg is “asdfasdfasdf”
- ▶ Of course, an attacker might not corrupt in this way...

A concrete stack overflow example

- ▶ Now, a stack snapshot with a malicious overflow:

<u>address</u>	<u>content</u>	
0x0012ff5c	0x00353037	; argument two pointer
0x0012ff58	0x0035302f	; argument one pointer
0x0012ff54	0x0012ff48	; return address: address of attack payload
0x0012ff50	0XXXXXXXXX	; irrelevant
0x0012ff4c	0XXXXXXXXX	; irrelevant
0x0012ff48	0feeb2ecd	; attack payload
0x0012ff44	0XX2f2f3a	; tmp continues ': ' '/' '/' ...
0x0012ff40	0x656c6966	; tmp array: 'f' 'i' 'l' 'e'



- ▶ In the above, the stack has been corrupted maliciously
- ▶ The args are “file:///” and particular attacker-chosen data
- ▶ XX can be any non-zero byte value

Stack canaries

- ▶ Very simple defense
 - ▶ Assume a contiguous buffer overflow is used by attackers
 - ▶ And that the overflow is based on zero-terminated strings
 - ▶ Put canary with “terminator” values below the return address

<u>address</u>	<u>content</u>	
0x0012ff5c	0x00353037	; argument two pointer
0x0012ff58	0x0035302f	; argument one pointer
0x0012ff54	0x00401263	; return address
0x0012ff50	0x0012ff7c	; saved base pointer
0x0012ff4c	0x00000000	; <i>all-zero canary</i>
0x0012ff48	0x00000072	; tmp continues 'r' '\0' '\0' '\0'
0x0012ff44	0x61626f6f	; tmp continues 'o' 'o' 'b' 'a'
0x0012ff40	0x662f2f3a	; tmp continues ':' '/' '/' 'f'
0x0012ff3c	0x656c6966	; tmp array: 'f' 'i' 'l' 'e'

- ▶ Check canary integrity before using return address!

Stack cookies

- ▶ Can also use random, secret values: **cookies**
- ▶ Defends against non-null-terminated overflows (e.g. via memcpy)

<u>address</u>	<u>content</u>	
0x0012ff5c	0x00353037	; argument two pointer
0x0012ff58	0x0035302f	; argument one pointer
0x0012ff54	0x00401263	; return address
0x0012ff50	0x0012ff7c	; saved base pointer
0x0012ff4c	0xF00DFEED	; a secret, random cookie value
0x0012ff48	0x00000072	; tmp continues 'r' '\0' '\0' '\0'
0x0012ff44	0x61626f6f	; tmp continues 'o' 'o' 'b' 'a'
0x0012ff40	0x662f2f3a	; tmp continues ':' '/' '/' 'f'
0x0012ff3c	0x656c6966	; tmp array: 'f' 'i' 'l' 'e'

- ▶ Check cookie integrity before using return address!
- ▶ Implemented in Windows (/GS compiler flag)

Stack canaries and cookies

- ▶ Stack canaries and stack cookies have very little cost
 - ▶ Only needed on functions with local arrays
 - ▶ Even so, not always applied: heuristics determine when
 - ▶ (Not a good idea, as shown by recent ANI attack on Vista)
- ▶ Widely implemented: /GS, StackGuard, ProPolice, etc.
 - ▶ Implementations typically combine with other defenses
- ▶ Main limitations:
 - ▶ Only protects against contiguous stack-based overflows
 - ▶ No protection if attack happens before function returns
 - ▶ For example, must protect function-pointer arguments
 - ▶ Do not prevent heap-based buffer overflows

Preventing data execution

- ▶ Simply prevent the execution of data as code
- ▶ This prevents both stack and heap-based attacks
- ▶ There is hardware support for this (NX bit on x86)
 - ▶ Can be done with pretty much zero overhead
 - ▶ But it breaks a lot of software:
 - ▶ Most Win32 GUI apps, CLR (and JITs)
- ▶ Can also be done in software (SMAC)
- ▶ Limitations:
 - ▶ Attackers don't always have to execute data as code
 - ▶ They can just corrupt data: data-only attacks
 - ▶ They can simply execute existing code: jump-to-libc

Jump-to-libc

- ▶ Any existing code can be executed by attackers
 - ▶ May be an existing function, such as `system()`
 - ▶ E.g., a function that is never invoked (dead code)
 - ▶ Or code in the middle of a function
- ▶ Can even be “opportunistic” code
 - ▶ Found within executable pages (e.g. switch tables)
 - ▶ Or found within existing instructions (long x86 instructions)
- ▶ Typically a step towards running attackers own shellcode
- ▶ These are *jump-to-libc* or *return-to-libc* attacks
- ▶ Allow attackers to overcome NX defenses

A new function to be attacked

- ▶ Computes the median integer in an input array
- ▶ Sorts a copy of the array and return the middle integer

```
int median( int* data, int len, void* cmp )
{
    // must have 0 < len <= MAX_INTS
    int tmp[MAX_INTS];
    memcpy( tmp, data, len*sizeof(int) );    // copy the input integers
    qsort( tmp, len, sizeof(int), cmp );    // sort the local copy
    return tmp[len/2];                      // median is in the middle
}
```

- ▶ If len is larger than MAX_INTS we have a stack overflow

An example bad function pointer

- ▶ Many ways to attack the median function
- ▶ The cmp pointer is used before the function returns
 - ▶ It can be overwritten by a stack-based overflow
 - ▶ And stack canaries or cookies are not a defense
- ▶ Using jump-to-libc, an attack can also foil NX
- ▶ Use existing code to install and jump to attack payload
 - ▶ Including marking the shellcode bytes as executable
- ▶ Example of *indirect code injection*
- ▶ (As opposed to *direct code injection* in previous attacks)

Concrete jump-to-libc attack example

- ▶ A normal stack for the median function
- ▶ Stack snapshot at the point of the call to memcpy
- ▶ MAX_INTS is 8
- ▶ The tmp array is empty, or all zero

stack address	normal stack contents
0x0012ff38	0x004013e0 ; cmp argument
0x0012ff34	0x00000001 ; len argument
0x0012ff30	0x00353050 ; data argument
0x0012ff2c	0x00401528 ; return address
0x0012ff28	0x0012ff4c ; saved base pointer
0x0012ff24	0x00000000 ; tmp final 4 bytes
0x0012ff20	0x00000000 ; tmp continues
0x0012ff1c	0x00000000 ; tmp continues
0x0012ff18	0x00000000 ; tmp continues
0x0012ff14	0x00000000 ; tmp continues
0x0012ff10	0x00000000 ; tmp continues
0x0012ff0c	0x00000000 ; tmp continues
0x0012ff08	0x00000000 ; tmp buffer starts
0x0012ff04	0x00000004 ; memcpy length argument
0x0012ff00	0x00353050 ; memcpy source argument
0x0012fefc	0x0012ff08 ; memcpy destination arg.

Concrete jump-to-libc attack example

- ▶ A **benign** stack overflow in the median function
- ▶ Not the values that an attacker will choose ...

stack address	benign overflow contents	
0x0012ff38	0x1111110d	; cmp argument
0x0012ff34	0x1111110c	; len argument
0x0012ff30	0x1111110b	; data argument
0x0012ff2c	0x1111110a	; return address
0x0012ff28	0x11111109	; saved base pointer
0x0012ff24	0x11111108	; tmp final 4 bytes
0x0012ff20	0x11111107	; tmp continues
0x0012ff1c	0x11111106	; tmp continues
0x0012ff18	0x11111105	; tmp continues
0x0012ff14	0x11111104	; tmp continues
0x0012ff10	0x11111103	; tmp continues
0x0012ff0c	0x11111102	; tmp continues
0x0012ff08	0x11111101	; tmp buffer starts
0x0012ff04	0x00000040	; memcpy length argument
0x0012ff00	0x00353050	; memcpy source argument
0x0012fefc	0x0012ff08	; memcpy destination arg.

Concrete jump-to-libc attack example

- ▶ A **malicious** stack overflow in the `median` function
- ▶ The attack doesn't use the return address (e.g., to avoid stack canary or cookie defenses)
- ▶ Control-flow is redirected in `qsort`
- ▶ Uses `jump-to-libc` to foil NX defenses

stack address	malicious overflow contents	
0x0012ff38	0x7c971649	; cmp argument
0x0012ff34	0x1111110c	; len argument
0x0012ff30	0x1111110b	; data argument
0x0012ff2c	0xfeeb2ecd	; return address
0x0012ff28	0x70000000	; saved base pointer
0x0012ff24	0x70000000	; tmp final 4 bytes
0x0012ff20	0x00000040	; tmp continues
0x0012ff1c	0x00003000	; tmp continues
0x0012ff18	0x00001000	; tmp continues
0x0012ff14	0x70000000	; tmp continues
0x0012ff10	0x7c80978e	; tmp continues
0x0012ff0c	0x7c809a51	; tmp continues
0x0012ff08	0x11111101	; tmp buffer starts
0x0012ff04	0x00000040	; memcpy length argument
0x0012ff00	0x00353050	; memcpy source argument
0x0012fefc	0x0012ff08	; memcpy destination arg.

Concrete jump-to-libc attack example

- ▶ Below shows the context of `cmp` invocation in `qsort`
- ▶ Goes to a 4-byte *trampoline* sequence found in a library

```
...  
push    edi                ; push second argument to be compared onto the stack  
push    ebx                ; push the first argument onto the stack  
call    [esp+comp_fp]      ; call comparison function, indirectly through a pointer  
add     esp, 8             ; remove the two arguments from the stack  
test    eax, eax           ; check the comparison result  
jle     label_lessthan     ; branch on that result
```

..

machine code				
address	opcode	bytes	assembly-language version of the machine code	
0x7c971649	0x8b	0xe3	<code>mov esp, ebx</code>	; change the stack location to ebx
0x7c97164b	0x5b		<code>pop ebx</code>	; pop ebx from the new stack
0x7c97164c	0xc3		<code>ret</code>	; return based on the new stack

The intent of the jump-to-libc attack

- ▶ Perform a series of calls to existing library functions
- ▶ With carefully selected arguments

```
// call a function to allocate writable, executable memory at 0x70000000  
VirtualAlloc(0x70000000, 0x1000, 0x3000, 0x40); // function at 0x7c809a51
```

```
// call a function to write the four-byte attack payload to 0x70000000  
InterlockedExchange(0x70000000, 0xfeeb2ecd); // function at 0x7c80978e
```

```
// invoke the four bytes of attack payload machine code  
((void (*)(void))0x70000000)(); // payload at 0x70000000
```

- ▶ The effect is to install and execute the attack payload

x86 `__cdecl` function-call convention

- ▶ Push parameters onto the stack, from right to left
- ▶ Call the function
- ▶ Save and update the `%ebp`
- ▶ Allocate local variables
- ▶ Perform the function's purpose
- ▶ Release local storage
- ▶ Restore saved registers
- ▶ Restore the old base pointer
- ▶ Return from the function
 - ▶ The `RET` instruction pops the old `%EIP` from the stack and jumps to that location. This gives control back to the caller function. Only the stack pointer and instruction pointers are modified by a subroutine return.
- ▶ Clean up pushed parameters

How the attack unwinds the stack

- ▶ First invalid control-flow edge goes to trampoline
- ▶ Trampoline returns to the start of VirtualAlloc
- ▶ Which returns to the start of the InterlockedExch. function
- ▶ Which returns to the copy of the attack payload

stack address	malicious overflow contents	
0x0012ff38	0x7c971649	; cmp argument
0x0012ff34	0x1111110c	; len
0x0012ff30	0x1111110b	; da
0x0012ff2c	0xfeeb2ecd	; re
0x0012ff28	0x70000000	; s
esp →	0x70000000	; tmp continues
0x0012ff20	0x00000040	; tmp
0x0012ff1c	0x00003000	; tmp
0x0012ff18	0x00001000	; tmp
0x0012ff14	0x70000000	; tmp continues
esp →	0x7c80978e	; tmp continues
0x0012ff0c	0x7c809a51	; tmp
0x0012ff08	0x11111101	; tmp
0x0012ff04	0x00000040	; memcpy length argument
0x0012ff00	0x00353050	; memcpy source argument
0x0012fefc	0x0012ff08	; memcpy destination arg.

New executable copy of attack payload

Interlocked Exchange

VirtualAlloc

Where to find useful trampolines?

- ▶ In Linux libc, one in 178 bytes is a 0xc3 ret opcode
- ▶ One in 475 bytes is an opportunistic, or unintended, ret

```
f7 c7 07 00 00 00    test    edi, 0x00000007
0f 95 45 c3          setnz   byte ptr [ebp-61]
```

Starting one byte later, the attacker instead obtains

```
c7 07 00 00 00 0f    movl    edi, 0x0f000000
95                   xchg    eax, ebp
45                   inc     ebp
c3                   ret
```

- ▶ All of these may be useful somehow

Generalized jump-to-libc attacks

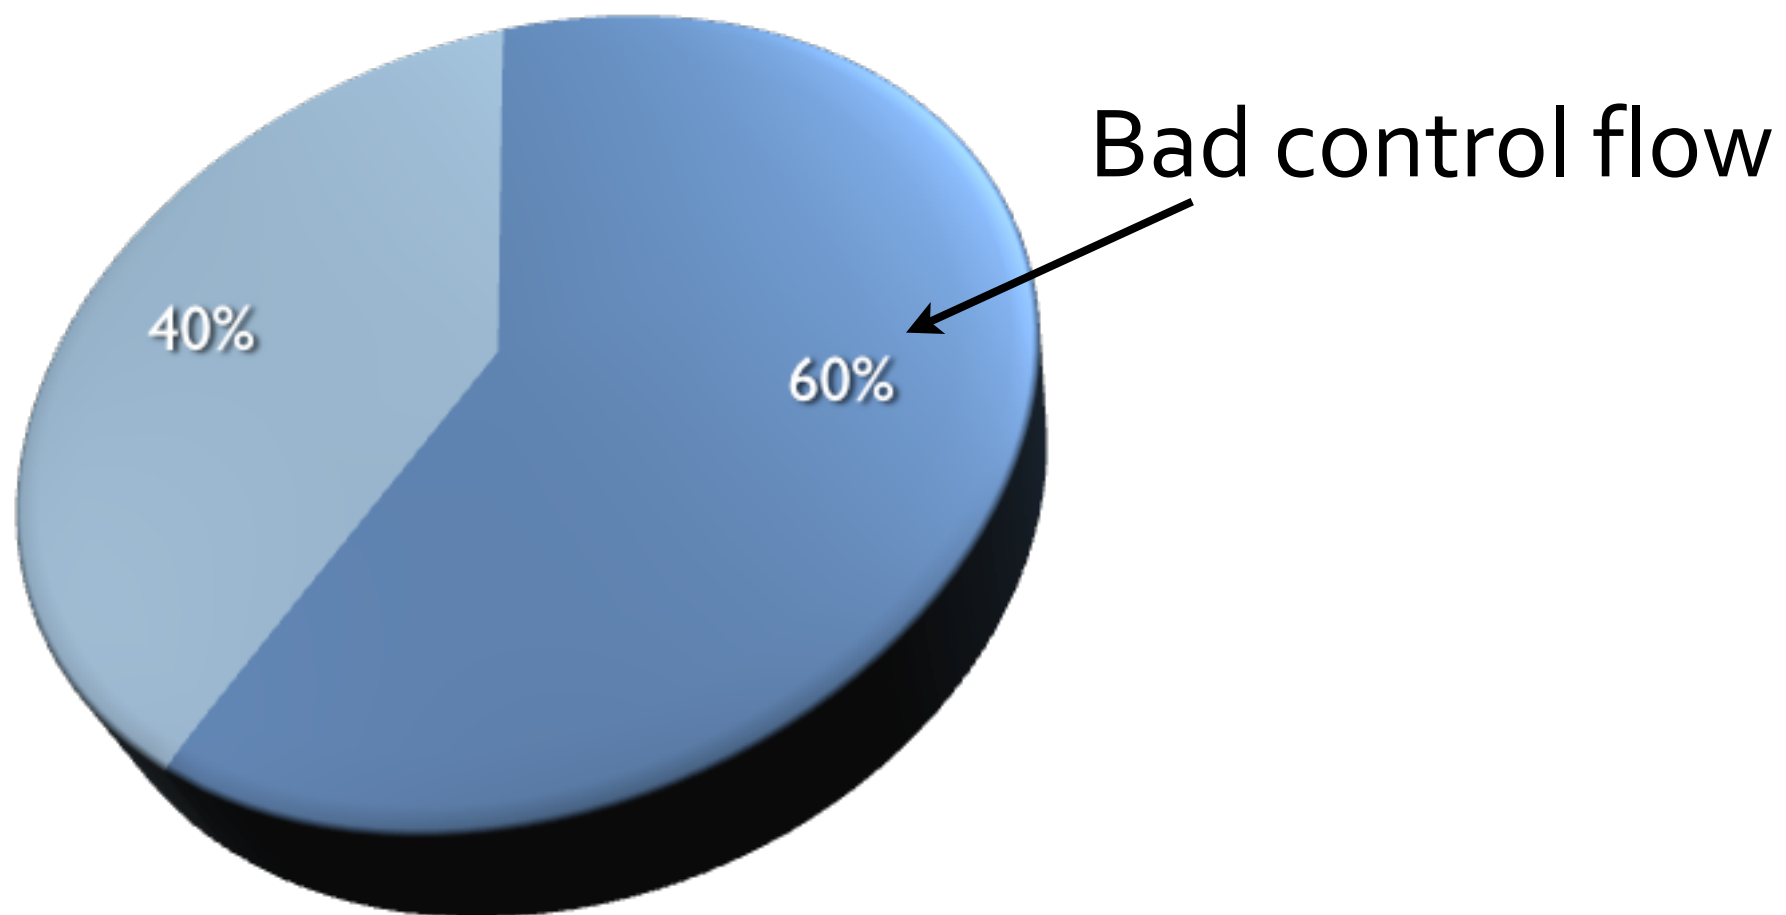
- ▶ Recent demonstration by Shacham [upcoming CCS'07]
 - ▶ Possible to achieve anything by only executing trampolines
 - ▶ Can compose trampolines into “gadget” primitives
 - ▶ Such “return-oriented-computing” is Turing complete
 - ▶ Practical, even if only opportunistic ret sequences are used
- ▶ Confirms a long-standing assumption:
*if arbitrary jumping around within existing,
executable code is permitted
then
an attacker can cause any desired, bad behavior*

Jump-to-libc attacks

- ▶ Jump-to-libc attacks are of great practical concern
 - ▶ For instance, recent ANI attack on Vista is similar to median
- ▶ Traditionally, return-to-libc with the target `system()`
 - ▶ Removing `system()` is neither a good nor sufficient defense
 - ▶ Generality of trampolines makes this a unarguable point
 - ▶ Anyway difficult to eliminate code from shared libraries
- ▶ Based on knowledge of existing code, and its addresses
 - ▶ Attackers must deal with natural software variability
 - ▶ Increasing the variability can be a good defense
- ▶ Best defense is to lock down the possible control flow
 - ▶ Other, simpler measures will also help

Control-flow integrity

- ▶ ~60% of attacks subvert the expected control flow
 - ▶ Enforcing control-flow integrity would prevent such attacks



Assumptions about control flow

- ▶ We write our code in high-level languages
- ▶ Naturally, our execution model assumes:
 - ▶ Functions start at the beginning
 - ▶ They (typically) execute from beginning to end
 - ▶ And, when done, they return to their call site
 - ▶ Only the code in the program can be executed
 - ▶ The set of executable instructions is limited to those output during compilation of the program

Assumptions about control flow

- ▶ We write our code in high-level languages
- ▶ **But, actually, at the level of machine code**
 - ▶ Can start in the middle of functions
 - ▶ A fragment of a function may be executed
 - ▶ Returns can go to any program instruction
 - ▶ All the data has usually been executable
 - ▶ On the x86, can start executing not only in the middle of functions, but middle of instructions!


What bytes will the CPU interpret?

- ▶ Hardware places few constraints on control flow
- ▶ A call to a function-pointer can lead many places:

Possible control flow destination

Safe code/data

Possible Execution of Memory



Enforcing control-flow integrity


- ▶ Only certain control-flow is possible in software
 - ▶ Even in C there are function and expression boundaries
 - ▶ Should also consider who-can-go-where, and dead code
- ▶ Control-flow integrity means that execution proceeds according to a specified control-flow graph (CFG).
 - ⇒ Reduces gap between machine code and high-level languages
- ▶ Can enforce with CFI mechanism, which is simple, efficient, and applicable to existing software.
- CFI enforces a basic property that thwarts a large class of attacks— without giving “end-to-end” security.

Guards for control flow integrity

- ▶ CFI guards restrict computed jumps and calls
 - ▶ Calls through function pointers (e.g. virtual methods in C++)
 - ▶ All return, exception and switch statements
- ▶ Direct calls are unaffected
- ▶ CFI guard matches label at source and target
 - ▶ Labels are constants embedded in machine-code
 - ▶ Labels are not secret, but must be **unique**
- ▶ Two destinations are equivalent when the CFG contains edges to it from the same set of sources
 - ▶ Equivalent destinations are labeled the same
 - ▶ i.e. a label uniquely identifies a CFG equivalence class

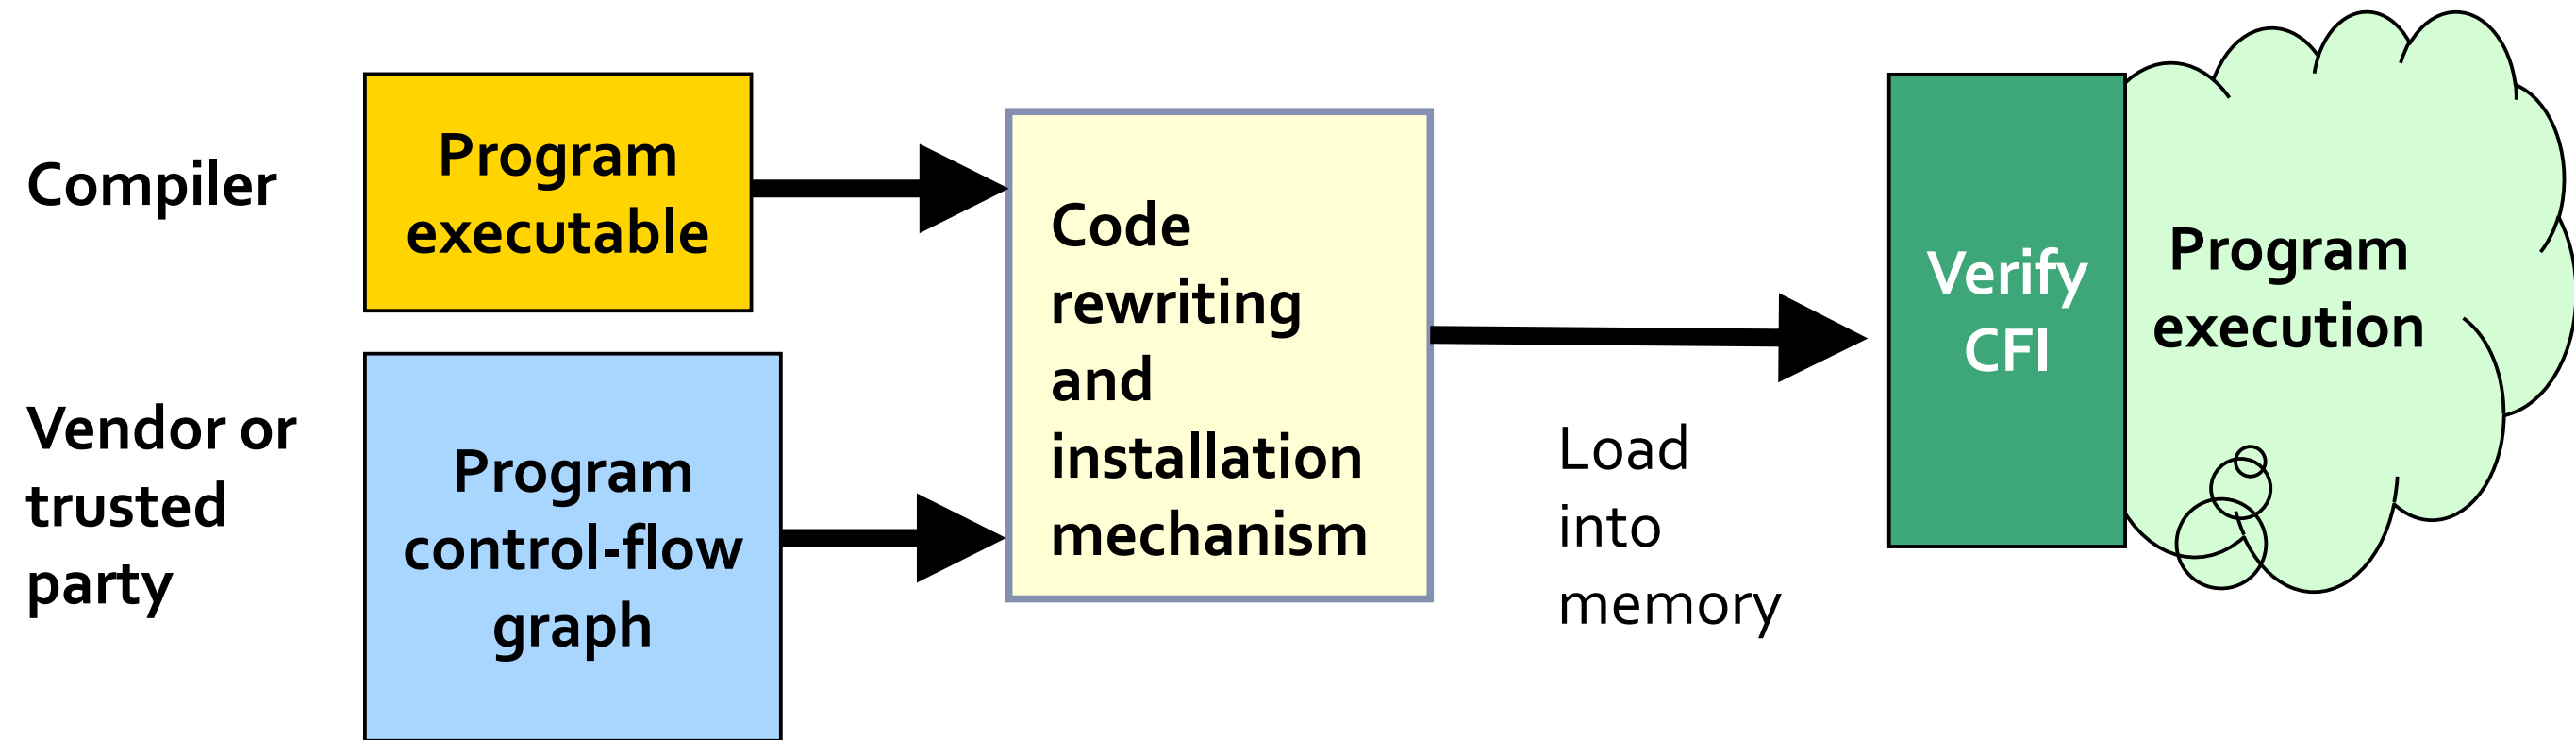
A simple example

```
bool lt(int x, int y) {  
    return x < y;  
}  
bool gt(int x, int y) {  
    return x > y;  
}  
  
sort2(int a[], int b[], int len)  
{  
    sort( a, len, lt );  
    sort( b, len, gt );  
}
```



- ▶ Ensure “labels” are correct at load- and run-time
 - ▶ Bit patterns identify different points in the code
 - ▶ Indirect control flow must go to the right pattern
- ▶ Can be enforced using software instrumentation
 - ▶ Even for existing, legacy software

Overview of a system with CFI



- ▶ Machine code rewriting using instrumentation tool
 - ▶ Applies to legacy Windows x86 executables
 - ▶ Code rewriting need not be trusted, because of the verifier
 - ▶ The verifier is simple (2 KLoC, mostly parsing x86 opcodes)


CFI formal study [ICFEM'05]

- ▶ Formally validated the benefits of CFI
 - ▶ Defined a machine code semantics
 - ▶ Powerful attacker model
 - ▶ Attacker can arbitrarily control all of data memory
 - ▶ Proved that, with CFI, execution always follows the CFG, even when under attack
- ▶ Assumptions
 - ▶ NXD: Data cannot be executed (hardware or software)
 - ▶ NWC: Code cannot be modified (hardware, already used)
 - ▶ We can rely on values in distinguished registers
 - ▶ Jumps cannot go into the middle of instructions
 - ▶ A convenient simplification to make the proof manageable

CFI as foundation for other prop.

- ▶ CFI can be used as a foundation for efficiently enforcing more sophisticated security properties
- ▶ Software fault isolation (e.g. sandboxing)
 - ▶ Dynamically check memory accesses to emulate traditional memory protection
- ▶ Software memory access control
 - ▶ Stronger than software fault isolation: isolated data memory regions accessible only from particular code
 - ▶ Removes NXD assumption, but adds extra overhead
- ▶ Protected shadow call stack
 - ▶ ID checks on return replaced by the use of a call stack
 - ▶ Very little extra overhead (at least with x86-specific tricks)

Cost and Benefits



- ▶ CFI overhead: ~16% in synthetic CPU-bound benchmarks
 - ▶ Is this really unobtrusive (close to zero overhead) ?
- ▶ Effectively stops most jump-to-libc attacks
 - ▶ No trampolining, even if CFI enforces a very coarse CFG
 - ▶ E.g., may have two labels -- for call sites and start of functions
- ▶ Limitation: Data-only attacks

Conclusion

- ▶ **Mitigation techniques**
 - ▶ Automatic defenses that work on legacy code
 - ▶ Operate at the machine-code level
 - ▶ Involve no source-code changes
 - ▶ Have close to zero overhead
 - ▶ Only prevent certain kinds of attacks
 - ▶ May provide a false feeling of security ... like a Volvo
 - ▶ Are not substitutes for correct code or safer languages
 - ▶ Do not protect against denial-of-service attacks
- ▶ **Control-flow integrity**
 - ▶ Particularly powerful mitigation technique
 - ▶ Prevents many kinds of attacks, including jump-to-libc